

[Dashboard](#) / [My courses](#) / [Computer Engineering & IT](#) / [CEIT-Even-sem-21-22](#) / [OS-even-sem-21-22](#) / [7 February - 13 February](#)
/ [Quiz-1: 10 AM](#)

Started on Saturday, 12 February 2022, 10:00:21 AM

State Finished

Completed on Saturday, 12 February 2022, 11:25:53 AM

Time taken 1 hour 25 mins

Grade 4.94 out of 10.00 (49%)

Question 1

Complete

Mark 0.00 out of 0.50

Select all the correct statements about code of bootmain() in xv6

```
void
bootmain(void)
{
    struct elfhdr *elf;
    struct proghdr *ph, *eph;
    void (*entry)(void);
    uchar* pa;

    elf = (struct elfhdr*)0x10000; // scratch space

    // Read 1st page off disk
    readseg((uchar*)elf, 4096, 0);

    // Is this an ELF executable?
    if(elf->magic != ELF_MAGIC)
        return; // let bootasm.S handle error

    // Load each program segment (ignores ph flags).
    ph = (struct proghdr*)((uchar*)elf + elf->phoff);
    eph = ph + elf->phnum;
    for(; ph < eph; ph++){
        pa = (uchar*)ph->paddr;
        readseg(pa, ph->filesz, ph->off);
        if(ph->memsz > ph->filesz)
            stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
    }

    // Call the entry point from the ELF header.
    // Does not return!
    entry = (void(*) (void))(elf->entry);
    entry();
}
```

Also, inspect the relevant parts of the xv6 code. binary files, etc and run commands as you deem fit to answer this question.

- ☒ a. The elf->entry is set by the linker in the kernel file and it's 0x80000000
- ☐ b. The condition if(ph->memsz > ph->filesz) is never true.
- ☒ c. The readseg finally invokes the disk I/O code using assembly instructions
- ☒ d. The elf->entry is set by the linker in the kernel file and it's 0x80000000
- ☐ e. The kernel ELF file contains actual physical address where particular sections of 'kernel' file should be loaded
- ☒ f. The kernel file gets loaded at the Physical address 0x10000 + 0x80000000 in memory.
- ☒ g. The kernel file in memory is not necessarily a continuously filled in chunk, it may have holes in it.
- ☐ h. The elf->entry is set by the linker in the kernel file and it's 8010000c
- ☐ i. The kernel file gets loaded at the Physical address 0x10000 in memory.
- ☐ j. The stosb() is used here, to fill in some space in memory with zeroes
- ☐ k. The kernel file has only two program headers

The correct answers are: The kernel file gets loaded at the Physical address 0x10000 in memory., The kernel file in memory is not necessarily a continuously filled in chunk, it may have holes in it., The elf->entry is set by the linker in the kernel file and it's 8010000c, The readseg finally invokes the disk I/O code using assembly instructions, The stosb() is used here, to fill in some space in memory with zeroes, The kernel ELF file contains actual physical address where particular sections of 'kernel' file should be loaded, The kernel file has only two program headers

Question 2

Complete

Mark 0.50 out of 0.50

What's the trapframe in xv6?

- ☐ a. A frame of memory that contains all the trap handler's addresses
- ☐ b. A frame of memory that contains all the trap handler code's function pointers
- ☐ c. The IDT table
- ☒ d. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware + code in trapasm.S
- ☐ e. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by code in trapasm.S only
- ☐ f. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware only
- ☐ g. A frame of memory that contains all the trap handler code

The correct answer is: The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware + code in trapasm.S

Question 3

Complete

Mark 0.21 out of 0.50

Order the events that occur on a timer interrupt:

Jump to a code pointed by IDT	3
Jump to scheduler code	2
Select another process for execution	5
Change to kernel stack of currently running process	4
Set the context of the new process	6
Save the context of the currently running process	1
Execute the code of the new process	7

The correct answer is: Jump to a code pointed by IDT → 2, Jump to scheduler code → 4, Select another process for execution → 5, Change to kernel stack of currently running process → 1, Set the context of the new process → 6, Save the context of the currently running process → 3, Execute the code of the new process → 7

Question 4

Complete

Mark 0.50 out of 0.50

Suppose a program does a `scanf()` call.

Essentially the `scanf` does a `read()` system call.

This call will obviously "block" waiting for the user input.

In terms of OS data structures and execution of code, what does it mean?

Select one:

- ☒ a. OS code for `read()` will move PCB of current process to a wait queue and call scheduler
- ☐ b. OS code for `read()` will call scheduler
- ☐ c. OS code for `read()` will move the PCB of this process to a wait queue and return from the system call
- ☐ d. `read()` will return and process will be taken to a wait queue
- ☐ e. `read()` returns and process calls scheduler()

The correct answer is: OS code for `read()` will move PCB of current process to a wait queue and call scheduler

Question 5

Complete

Mark 0.50 out of 0.50

In `bootasm.S`, on the line

```
ljmp    $(SEG_KCODE<<3), $start32
```

The `SEG_KCODE << 3`, that is shifting of 1 by 3 bits is done because

- ☐ a. The value 8 is stored in code segment
- ☒ b. The code segment is 16 bit and only upper 13 bits are used for segment number
- ☐ c. While indexing the GDT using CS, the value in CS is always divided by 8
- ☐ d. The `ljmp` instruction does a divide by 8 on the first argument
- ☐ e. The code segment is 16 bit and only lower 13 bits are used for segment number

The correct answer is: The code segment is 16 bit and only upper 13 bits are used for segment number

Question 6

Complete

Mark 0.40 out of 0.50

Select Yes if the mentioned element should be a part of PCB

Select No otherwise.

Yes	No	
<input checked="" type="radio"/>	<input type="radio"/>	Memory management information about that process
<input checked="" type="radio"/>	<input type="radio"/>	Process context
<input type="radio"/>	<input checked="" type="radio"/>	Function pointers to all system calls
<input checked="" type="radio"/>	<input type="radio"/>	PID
<input checked="" type="radio"/>	<input type="radio"/>	EIP at the time of context switch
<input checked="" type="radio"/>	<input type="radio"/>	List of opened files
<input checked="" type="radio"/>	<input type="radio"/>	PID of Init
<input checked="" type="radio"/>	<input type="radio"/>	Process state
<input type="radio"/>	<input checked="" type="radio"/>	Pointer to the parent process
<input type="radio"/>	<input checked="" type="radio"/>	Pointer to IDT

Memory management information about that process: Yes

Process context: Yes

Function pointers to all system calls: No

PID: Yes

EIP at the time of context switch: Yes

List of opened files: Yes

PID of Init: No

Process state: Yes

Pointer to the parent process: Yes

Pointer to IDT: No

Question 7

Complete

Mark 0.40 out of 1.00

Mark the statements, w.r.t. the scheduler of xv6 as True or False

True	False	
<input type="radio"/>	<input checked="" type="radio"/>	The work of selecting and scheduling a process is done only in <code>scheduler()</code> and not in <code>sched()</code>
<input checked="" type="radio"/>	<input type="radio"/>	<code>swtch</code> is a function that saves old context, loads new context, and returns to last EIP in the new context
<input type="radio"/>	<input checked="" type="radio"/>	The variable <code>c->scheduler</code> on first processor uses the stack allocated entry.S
<input checked="" type="radio"/>	<input type="radio"/>	<code>sched()</code> and <code>scheduler()</code> are co-routines
<input type="radio"/>	<input checked="" type="radio"/>	The function <code>scheduler()</code> executes using the kernel-only stack
<input type="radio"/>	<input checked="" type="radio"/>	When a process is scheduled for execution, it resumes execution in <code>sched()</code> after the call to <code>swtch()</code>
<input type="radio"/>	<input checked="" type="radio"/>	<code>swtch</code> is a function that does not return to the caller
<input type="radio"/>	<input checked="" type="radio"/>	the control returns to <code>mycpu()->intena = intena; ();</code> after <code>swtch(&p->context, mycpu()->scheduler);</code> in <code>sched()</code>
<input type="radio"/>	<input checked="" type="radio"/>	the control returns to <code>switchkvm();</code> after <code>swtch(&(c->scheduler), p->context);</code> in <code>scheduler()</code>
<input checked="" type="radio"/>	<input type="radio"/>	<code>sched()</code> calls <code>scheduler()</code> and <code>scheduler()</code> calls <code>sched()</code>

The work of selecting and scheduling a process is done only in `scheduler()` and not in `sched()`: True`swtch` is a function that saves old context, loads new context, and returns to last EIP in the new context: TrueThe variable `c->scheduler` on first processor uses the stack allocated entry.S: True`sched()` and `scheduler()` are co-routines: TrueThe function `scheduler()` executes using the kernel-only stack: TrueWhen a process is scheduled for execution, it resumes execution in `sched()` after the call to `swtch()`

: True

`swtch` is a function that does not return to the caller: Truethe control returns to `mycpu()->intena = intena; ();` after `swtch(&p->context, mycpu()->scheduler);` in `sched()`:

False

the control returns to `switchkvm();` after `swtch(&(c->scheduler), p->context);` in `scheduler()`: False`sched()` calls `scheduler()` and `scheduler()` calls `sched()`: False

Question 8

Complete

Mark 0.00 out of 0.50

Consider the following programs

exec1.c

```
#include <unistd.h>
#include <stdio.h>
int main() {
    execl("./exec2", "./exec2", NULL);
}
```

exec2.c

```
#include <unistd.h>
#include <stdio.h>
int main() {
    execl("/bin/ls", "/bin/ls", NULL);

    printf("hello\n");
}
```

Compiled as

```
cc  exec1.c -o exec1
cc  exec2.c -o exec2
```

And run as

```
$ ./exec1
```

Explain the output of the above command (./exec1)

Assume that /bin/ls , i.e. the 'ls' program exists.

Select one:

- ☐ a. "ls" runs on current directory
- ☐ b. Execution fails as the call to execl() in exec1 fails
- ☐ c. Execution fails as one exec can't invoke another exec
- ☒ d. Program prints hello
- ☐ e. Execution fails as the call to execl() in exec2 fails

The correct answer is: "ls" runs on current directory

Question 9

Complete

Mark 0.00 out of 0.50

For each line of code mentioned on the left side, select the location of sp/esp that is in use

cli

in bootasm.S

0x7c00 to 0

readseg((uchar*)elf, 4096, 0);

in bootmain.c

Immaterial as the stack is not used here

ljmp \$(SEG_KCODE<<3), \$start32

in bootasm.S

0x10000 to 0x7c00

call bootmain

in bootasm.S

The 4KB area in kernel image, loaded in memory, named as 'stack'

jmp *%eax

in entry.S

0x7c00 to 0x10000

The correct answer is: cli

in bootasm.S → Immaterial as the stack is not used here, readseg((uchar*)elf, 4096, 0);

in bootmain.c → 0x7c00 to 0, ljmp \$(SEG_KCODE<<3), \$start32

in bootasm.S → Immaterial as the stack is not used here, call bootmain

in bootasm.S → 0x7c00 to 0, jmp *%eax

in entry.S → The 4KB area in kernel image, loaded in memory, named as 'stack'

Question **10**

Complete

Mark 0.63 out of 1.00

Select the correct statements about interrupt handling in xv6 code

- ☒ a. xv6 uses the 64th entry in IDT for system calls
- ☐ b. xv6 uses the 0x64th entry in IDT for system calls
- ☒ c. On any interrupt/syscall/exception the control first jumps in vectors.S
- ☐ d. The function trap() is the called irrespective of hardware interrupt/system-call/exception
- ☒ e. Before going to alltraps, the kernel stack contains upto 5 entries.
- ☐ f. Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt
- ☒ g. The trapframe pointer in struct proc, points to a location on kernel stack
- ☐ h. The function trap() is the called only in case of hardware interrupt
- ☐ i. The trapframe pointer in struct proc, points to a location on user stack
- ☐ j. On any interrupt/syscall/exception the control first jumps in trapasm.S
- ☒ k. The CS and EIP are changed only after pushing user code's SS,ESP on stack
- ☐ l. The CS and EIP are changed only immediately on a hardware interrupt
- ☐ m. All the 256 entries in the IDT are filled

The correct answers are: All the 256 entries in the IDT are filled, Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt, xv6 uses the 64th entry in IDT for system calls, On any interrupt/syscall/exception the control first jumps in vectors.S, Before going to alltraps, the kernel stack contains upto 5 entries., The trapframe pointer in struct proc, points to a location on kernel stack, The function trap() is the called irrespective of hardware interrupt/system-call/exception, The CS and EIP are changed only after pushing user code's SS,ESP on stack

Question **11**

Complete

Mark 0.50 out of 0.50

Order the sequence of events, in scheduling process P1 after process P0

timer interrupt occurs

Process P0 is running

context of P1 is loaded from P1's PCB

Process P1 is running

Control is passed to P1

context of P0 is saved in P0's PCB

The correct answer is: timer interrupt occurs → 2, Process P0 is running → 1, context of P1 is loaded from P1's PCB → 4, Process P1 is running → 6, Control is passed to P1 → 5, context of P0 is saved in P0's PCB → 3

Question **12**

Complete

Mark 0.00 out of 1.00

Select the sequence of events that are NOT possible, assuming a non-interruptible kernel code

(Note: non-interruptible kernel code means, if the kernel code is executing, then interrupts will be disabled).

Note: A possible sequence may have some missing steps in between. An impossible sequence will have n and $n+1$ th steps such that $n+1$ th step can not follow n 'th step.

Select one or more:

- ☐ a. P1 running
P1 makes system call
timer interrupt
Scheduler
P2 running
timer interrupt
Scheduler
P1 running
P1's system call return
- ☐ b. P1 running
P1 makes system call
system call returns
P1 running
timer interrupt
Scheduler running
P2 running
- ☐ c. P1 running
P1 makes system call and blocks
Scheduler
P2 running
P2 makes system call and blocks
Scheduler
P3 running
Hardware interrupt
Interrupt unblocks P1
Interrupt returns
P3 running
Timer interrupt
Scheduler
P1 running
- ☒ d. P1 running
keyboard hardware interrupt
keyboard interrupt handler running
interrupt handler returns
P1 running
P1 makes system call
system call returns
P1 running
timer interrupt
scheduler
P2 running
- ☐ e.
P1 running
P1 makes system call
Scheduler
P2 running
P2 makes system call and blocks
Scheduler
P1 running again

- ☐ f. P1 running
 P1 makes sytem call and blocks
 Scheduler
 P2 running
 P2 makes sytem call and blocks
 Scheduler
 P1 running again

The correct answers are: P1 running
 P1 makes sytem call and blocks
 Scheduler
 P2 running
 P2 makes sytem call and blocks
 Scheduler
 P1 running again, P1 running
 P1 makes system call
 timer interrupt
 Scheduler
 P2 running
 timer interrupt
 Scheuler
 P1 running
 P1's system call return,
 P1 running
 P1 makes sytem call
 Scheduler
 P2 running
 P2 makes sytem call and blocks
 Scheduler
 P1 running again

Question 13

Complete

Mark 0.50 out of 0.50

Some part of the bootloader of xv6 is written in assembly while some part is written in C. Why is that so?
 Select all the appropriate choices

- ☒ a. The code in assembly is required for transition to protected mode, from real mode; after that calling convention applies, hence code can be written in C
- ☒ b. The setting up of the most essential memory management infrastructure needs assembly code
- ☐ c. The code in assembly is required for transition to protected mode, from real mode; but calling convention was applicable all the time
- ☐ d. The code for reading ELF file can not be written in assembly

The correct answers are: The code in assembly is required for transition to protected mode, from real mode; after that calling convention applies, hence code can be written in C, The setting up of the most essential memory management infrastructure needs assembly code

Question **14**

Complete

Mark 0.17 out of 0.50

The bootmain() function has this code

```
elf = (struct elfhdr*)0x10000; // scratch space
readseg((uchar*)elf, 4096, 0);
```

Mark the statements as True or False with respect to this code.

In these statements 0x1000 is referred to as ADDRESS

True	False	
<input checked="" type="radio"/>	<input type="radio"/>	The value ADDRESS is changed to a 0 the program could still work
<input checked="" type="radio"/>	<input type="radio"/>	It the value of ADDRESS is changed to a higher number (upto a limit), the program could still work
<input checked="" type="radio"/>	<input type="radio"/>	This line loads the kernel code at ADDRESS
<input checked="" type="radio"/>	<input type="radio"/>	It the value of ADDRESS is changed to a lower number (upto a limit), the program could still work
<input checked="" type="radio"/>	<input type="radio"/>	This line effectively loads the ELF header and the program headers at ADDRESS
<input checked="" type="radio"/>	<input type="radio"/>	If the value of ADDRESS is changed, then the program will not work

The value ADDRESS is changed to a 0 the program could still work: False

It the value of ADDRESS is changed to a higher number (upto a limit), the program could still work: True

This line loads the kernel code at ADDRESS: False

It the value of ADDRESS is changed to a lower number (upto a limit), the program could still work: True

This line effectively loads the ELF header and the program headers at ADDRESS: False

If the value of ADDRESS is changed, then the program will not work: False

Question **15**

Complete

Mark 0.50 out of 1.00

Which parts of the xv6 code in bootasm.S bootmain.c , entry.S and in the codepath related to scheduler() and trap handling() can also be written in some other way, and still ensure that xv6 works properly?

Writing code is not necessary. You only need to comment on which part of the code could be changed to something else or written in another fashion.

Maximum two points to be written.

We can use a scheduling algorithm. We can use the kernel stack in scheduler function in entry.S and bootmain.c .

Question **16**

Complete

Mark 0.13 out of 0.50

Select all the correct statements about zombie processes

Select one or more:

- ☐ a. If the parent of a process finishes, before the process itself, then after finishing the process is typically attached to 'init' as parent
- ☐ b. Zombie processes are harmless even if OS is up for long time
- ☒ c. A zombie process occupies space in OS data structures
- ☐ d. A process becomes zombie when it finishes, and remains zombie until parent calls wait() on it
- ☒ e. A process can become zombie if it finishes, but the parent has finished before it
- ☒ f. init() typically keeps calling wait() for zombie processes to get cleaned up
- ☐ g. A process becomes zombie when it's parent finishes
- ☒ h. A zombie process remains zombie forever, as there is no way to clean it up

The correct answers are: A process becomes zombie when it finishes, and remains zombie until parent calls wait() on it, A process can become zombie if it finishes, but the parent has finished before it, A zombie process occupies space in OS data structures, If the parent of a process finishes, before the process itself, then after finishing the process is typically attached to 'init' as parent, init() typically keeps calling wait() for zombie processes to get cleaned up

◀ [Extra Reading on Linkers: A writeup by Ian Taylor](#) (keep changing url string from 38 to 39, and so on)

Jump to...

[Dashboard](#) / [My courses](#) / [Computer Engineering & IT](#) / [CEIT-Even-sem-20-21](#) / [OS-Even-sem-2020-21](#) / 14 February - 20 February
/ [Quiz-1](#)

Started on Saturday, 20 February 2021, 2:51 PM

State Finished

Completed on Saturday, 20 February 2021, 3:55 PM

Time taken 1 hour 3 mins

Grade 7.30 out of 20.00 (37%)

Question 1

Partially correct

Mark 0.80 out of 1.00

Select all the correct statements about the state of a process.

- ☒ a. A process can self-terminate only when it's running ✓
- ☒ b. Typically, it's represented as a number in the PCB ✓
- ☒ c. A process that is running is not on the ready queue ✓
- ☒ d. Processes in the ready queue are in the ready state ✓
- ☐ e. It is not maintained in the data structures by kernel, it is only for conceptual understanding of programmers
- ☒ f. Changing from running state to waiting state results in "giving up the CPU" ✓
- ☐ g. A process in ready state is ready to receive interrupts
- ☒ h. A waiting process starts running after the wait is over ✗
- ☒ i. A process changes from running to ready state on a timer interrupt ✓
- ☒ j. A process in ready state is ready to be scheduled ✓
- ☒ k. A running process may terminate, or go to wait or become ready again ✓
- ☒ l. A process waiting for I/O completion is typically woken up by the particular interrupt handler code ✓
- ☐ m. A process waiting for any condition is woken up by another process only
- ☐ n. A process changes from running to ready state on a timer interrupt or any I/O wait

Your answer is partially correct.

You have selected too many options.

The correct answers are: Typically, it's represented as a number in the PCB, A process in ready state is ready to be scheduled, Processes in the ready queue are in the ready state, A process that is running is not on the ready queue, A running process may terminate, or go to wait or become ready again, A process changes from running to ready state on a timer interrupt, Changing from running state to waiting state results in "giving up the CPU", A process can self-terminate only when it's running, A process waiting for I/O completion is typically woken up by the particular interrupt handler code

Question 2

Incorrect

Mark 0.00 out of 1.00

For each line of code mentioned on the left side, select the location of sp/esp that is in use

<code>jmp *%eax</code> in <code>entry.S</code>	0x7c00 to 0x10000	✗
<code>ljmp \$(SEG_KCODE<<3), \$start32</code> in <code>bootasm.S</code>	0x10000 to 0x7c00	✗
<code>call bootmain</code> in <code>bootasm.S</code>	0x7c00 to 0x10000	✗
<code>cli</code> in <code>bootasm.S</code>	0x7c00 to 0	✗
<code>readseg((uchar*)elf, 4096, 0);</code> in <code>bootmain.c</code>	The 4KB area in kernel image, loaded in memory, named as 'stack'	✗

Your answer is incorrect.

The correct answer is: `jmp *%eax`

in `entry.S` → The 4KB area in kernel image, loaded in memory, named as 'stack', `ljmp $(SEG_KCODE<<3), $start32`

in `bootasm.S` → Immaterial as the stack is not used here, `call bootmain`

in `bootasm.S` → 0x7c00 to 0, `cli`

in `bootasm.S` → Immaterial as the stack is not used here, `readseg((uchar*)elf, 4096, 0);`

in `bootmain.c` → 0x7c00 to 0

Question 3

Correct

Mark 0.25 out of 0.25

Order the following events in boot process (from 1 onwards)

Boot loader	2	✓
Shell	6	✓
BIOS	1	✓
OS	3	✓
Init	4	✓
Login interface	5	✓

Your answer is correct.

The correct answer is: Boot loader → 2, Shell → 6, BIOS → 1, OS → 3, Init → 4, Login interface → 5

Question 4

Partially correct

Mark 0.30 out of 0.50

Consider the following command and its output:

```
$ ls -lht xv6.img kernel
-rw-rw-r-- 1 abhijit abhijit 4.9M Feb 15 11:09 xv6.img
-rwxrwxr-x 1 abhijit abhijit 209K Feb 15 11:09 kernel*
```

Following code in bootmain()

```
readseg((uchar*)elf, 4096, 0);
```

and following selected lines from Makefile

```
xv6.img: bootblock kernel
    dd if=/dev/zero of=xv6.img count=10000
    dd if=bootblock of=xv6.img conv=notrunc
    dd if=kernel of=xv6.img seek=1 conv=notrunc
```

```
kernel: $(OBJJS) entry.o entryother initcode kernel.ld
    $(LD) $(LDFLAGS) -T kernel.ld -o kernel entry.o $(OBJJS) -b binary initcode entryother
    $(OBJDUMP) -S kernel > kernel.asm
    $(OBJDUMP) -t kernel | sed '1,/SYMBOL TABLE/d; s/ .* / /; /^$$/d' > kernel.sym
```

Also read the code of bootmain() in xv6 kernel.

Select the options that describe the meaning of these lines and their correlation.

- ☐ a. Although the size of the kernel file is 209 Kb, only 4Kb out of it is the actual kernel code and remaining part is all zeroes.
- ☒ b. The kernel is compiled by linking multiple .o files created from .c files; and the entry.o, initcode, entryother files ✓
- ☒ c. The kernel.ld file contains instructions to the linker to link the kernel properly ✓
- ☐ d. The bootmain() code does not read the kernel completely in memory
- ☐ e. readseg() reads first 4k bytes of kernel in memory
- ☐ f. Although the size of the xv6.img file is ~5MB, only some part out of it is the bootloader+kernel code and remaining part is all zeroes.
- ☐ g. The kernel.asm file is the final kernel file
- ☐ h. The kernel disk image is ~5MB, the kernel within it is 209 kb, but bootmain() initially reads only first 4kb, and the later part is not read as it is user programs.
- ☒ i. The kernel disk image is ~5MB, the kernel within it is 209 kb, but bootmain() initially reads only first 4kb, and the later part is read using program headers in bootmain(). ✓

Your answer is partially correct.

You have correctly selected 3.

The correct answers are: The kernel disk image is ~5MB, the kernel within it is 209 kb, but bootmain() initially reads only first 4kb, and the later part is read using program headers in bootmain()., readseg() reads first 4k bytes of kernel in memory, The kernel is compiled by linking multiple .o files created from .c files; and the entry.o, initcode, entryother files, The kernel.ld file contains instructions to the linker to link the kernel properly, Although the size of the xv6.img file is ~5MB, only some part out of it is the bootloader+kernel code and remaining part is all zeroes.

Question **5**

Partially correct

Mark 0.50 out of 1.00

```
int f() {  
    int count;  
    for (count = 0; count < 2; count++) {  
        if (fork() == 0)  
            printf("Operating-System\n");  
    }  
    printf("TYCOMP\n");  
}
```

The number of times "Operating-System" is printed, is:

Answer: ☒

The correct answer is: 7.00

Question 6

Partially correct

Mark 0.40 out of 0.50

Select Yes/True if the mentioned element must be a part of PCB

Select No/False otherwise.

Yes	No		
<input checked="" type="radio"/>	<input type="radio"/>	PID	✓
<input checked="" type="radio"/>	<input type="radio"/>	Process context	✓
<input checked="" type="radio"/>	<input type="radio"/>	List of opened files	✓
<input checked="" type="radio"/>	<input type="radio"/>	Process state	✓
<input type="radio"/>	<input checked="" type="radio"/>	Parent's PID	✗
<input type="radio"/>	<input checked="" type="radio"/>	Pointer to IDT	✓
<input type="radio"/>	<input checked="" type="radio"/>	Function pointers to all system calls	✓
<input checked="" type="radio"/>	<input type="radio"/>	Memory management information about that process	✓
<input checked="" type="radio"/>	<input type="radio"/>	Pointer to the parent process	✗
<input checked="" type="radio"/>	<input type="radio"/>	EIP at the time of context switch	✓

PID: Yes

Process context: Yes

List of opened files: Yes

Process state: Yes

Parent's PID: No

Pointer to IDT: No

Function pointers to all system calls: No

Memory management information about that process: Yes

Pointer to the parent process: Yes

EIP at the time of context switch: Yes

Question 7

Incorrect

Mark 0.00 out of 1.00

Select all the correct statements about code of bootmain() in xv6

```
void
bootmain(void)
{
    struct elfhdr *elf;
    struct proghdr *ph, *eph;
    void (*entry)(void);
    uchar* pa;

    elf = (struct elfhdr*)0x10000; // scratch space

    // Read 1st page off disk
    readseg((uchar*)elf, 4096, 0);

    // Is this an ELF executable?
    if(elf->magic != ELF_MAGIC)
        return; // let bootasm.S handle error

    // Load each program segment (ignores ph flags).
    ph = (struct proghdr*)((uchar*)elf + elf->phoff);
    eph = ph + elf->phnum;
    for(; ph < eph; ph++){
        pa = (uchar*)ph->paddr;
        readseg(pa, ph->filesz, ph->off);
        if(ph->memsz > ph->filesz)
            stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
    }

    // Call the entry point from the ELF header.
    // Does not return!
    entry = (void(*) (void))(elf->entry);
    entry();
}
```

Also, inspect the relevant parts of the xv6 code. binary files, etc and run commands as you deem fit to answer this question.

- ☒ a. The kernel file gets loaded at the Physical address 0x10000 + 0x80000000 in memory. ✗
- ☒ b. The elf->entry is set by the linker in the kernel file and it's 0x80000000 ✗
- ☒ c. The kernel ELF file contains actual physical address where particular sections of 'kernel' file should be loaded ✓
- ☒ d. The kernel file in memory is not necessarily a continuously filled in chunk, it may have holes in it. ✓
- ☒ e. The kernel file has only two program headers ✓
- ☒ f. The elf->entry is set by the linker in the kernel file and it's 0x80000000 ✗
- ☒ g. The readseg finally invokes the disk I/O code using assembly instructions ✓
- ☒ h. The elf->entry is set by the linker in the kernel file and it's 8010000c ✓
- ☒ i. The kernel file gets loaded at the Physical address 0x10000 in memory. ✓
- ☒ j. The condition if(ph->memsz > ph->filesz) is never true. ✗
- ☒ k. The stosb() is used here, to fill in some space in memory with zeroes ✓

Your answer is incorrect.

The correct answers are: The kernel file gets loaded at the Physical address 0x10000 in memory., The kernel file in memory is not necessarily a continuously filled in chunk, it may have holes in it., The elf->entry is set by the linker in the kernel file and it's 8010000c, The readseg finally invokes the disk I/O code using assembly instructions, The stosb() is used here, to fill in some space in memory with zeroes, The kernel ELF file contains actual physical address where particular sections of 'kernel' file should be loaded, The kernel file has only two program headers

Question 8

Partially correct

Mark 0.13 out of 0.25

Which of the following are NOT a part of job of a typical compiler?

- ☒ a. Check the program for logical errors
- ☐ b. Convert high level language code to machine code
- ☐ c. Process the # directives in a C program
- ☐ d. Invoke the linker to link the function calls with their code, extern globals with their declaration
- ☐ e. Check the program for syntactical errors
- ☐ f. Suggest alternative pieces of code that can be written



Your answer is partially correct.

You have correctly selected 1.

The correct answers are: Check the program for logical errors, Suggest alternative pieces of code that can be written

Question 9

Correct

Mark 0.25 out of 0.25

Rank the following storage systems from slowest (first) to fastest(last)

Cache	6	✓
Hard Disk	3	✓
RAM	5	✓
Optical Disks	2	✓
Non volatile memory	4	✓
Registers	7	✓
Magnetic Tapes	1	✓

Your answer is correct.

The correct answer is: Cache → 6, Hard Disk → 3, RAM → 5, Optical Disks → 2, Non volatile memory → 4, Registers → 7, Magnetic Tapes → 1

Question **10**

Partially correct

Mark 0.21 out of 0.50

Which of the following parts of a C program do not have any corresponding machine code ?

- ☐ a. local variable declaration
- ☐ b. global variables
- ☒ c. function calls
- ☒ d. #directives
- ☐ e. expressions
- ☐ f. pointer dereference
- ☒ g. typedefs

✗

✓

✓

Your answer is partially correct.

You have correctly selected 2.

The correct answers are: #directives, typedefs, global variables

Question **11**

Correct

Mark 0.25 out of 0.25

Match a system call with it's description

pipe	create an unnamed FIFO storage with 2 ends - one for reading and another for writing	✓
dup	create a copy of the specified file descriptor into smallest available file descriptor	✓
dup2	create a copy of the specified file descriptor into another specified file descriptor	✓
exec	execute a binary file overlaying the image of current process	✓
fork	create an identical child process	✓

Your answer is correct.

The correct answer is: pipe → create an unnamed FIFO storage with 2 ends - one for reading and another for writing, dup → create a copy of the specified file descriptor into smallest available file descriptor, dup2 → create a copy of the specified file descriptor into another specified file descriptor, exec → execute a binary file overlaying the image of current process, fork → create an identical child process

Question **12**

Correct

Mark 0.25 out of 0.25

Match the register with the segment used with it.

eip	cs	✓
edi	es	✓
esi	ds	✓
ebp	ss	✓
esp	ss	✓

Your answer is correct.

The correct answer is: eip → cs, edi → es, esi → ds, ebp → ss, esp → ss

Question **13**

Correct

Mark 0.25 out of 0.25

What's the trapframe in xv6?

- ☐ a. A frame of memory that contains all the trap handler code
- ☐ b. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware only
- ☐ c. The IDT table
- ☐ d. A frame of memory that contains all the trap handler code's function pointers
- ☐ e. A frame of memory that contains all the trap handler's addresses
- ☒ f. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware + code in trapasm.S ✓
- ☐ g. The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by code in trapasm.S only

Your answer is correct.

The correct answer is: The sequence of values, including saved registers, constructed on the stack when an interrupt occurs, built by hardware + code in trapasm.S

Question 14

Incorrect

Mark 0.00 out of 0.50

Select all the correct statements about linking and loading.

Select one or more:

- ☒ a. Continuous memory management schemes can support dynamic linking and dynamic loading. ✗
- ☒ b. Loader is last stage of the linker program ✗
- ☒ c. Continuous memory management schemes can support static linking and dynamic loading. (may be inefficiently) ✓
- ☒ d. Dynamic linking and loading is not possible without demand paging or demand segmentation. ✓
- ☒ e. Dynamic linking essentially results in relocatable code. ✓
- ☒ f. Continuous memory management schemes can support static linking and static loading. (may be inefficiently) ✓
- ☒ g. Loader is part of the operating system ✓
- ☒ h. Static linking leads to non-relocatable code ✗
- ☒ i. Dynamic linking is possible with continuous memory management, but variable sized partitions only. ✗

Your answer is incorrect.

The correct answers are: Continuous memory management schemes can support static linking and static loading. (may be inefficiently), Continuous memory management schemes can support static linking and dynamic loading. (may be inefficiently), Dynamic linking essentially results in relocatable code., Loader is part of the operating system, Dynamic linking and loading is not possible without demand paging or demand segmentation.

Question 15

Incorrect

Mark 0.00 out of 0.25

In bootasm.S, on the line

```
ljmp $(SEG_KCODE<<3), $start32
```

The SEG_KCODE << 3, that is shifting of 1 by 3 bits is done because

- ☐ a. The value 8 is stored in code segment
- ☐ b. The code segment is 16 bit and only upper 13 bits are used for segment number
- ☒ c. The code segment is 16 bit and only lower 13 bits are used for segment number ✗
- ☐ d. While indexing the GDT using CS, the value in CS is always divided by 8
- ☐ e. The ljmp instruction does a divide by 8 on the first argument

Your answer is incorrect.

The correct answer is: The code segment is 16 bit and only upper 13 bits are used for segment number

Question **16**

Partially correct

Mark 0.07 out of 0.50

Order the events that occur on a timer interrupt:

Change to kernel stack	1	✗
Jump to a code pointed by IDT	2	✗
Jump to scheduler code	5	✗
Set the context of the new process	4	✗
Save the context of the currently running process	3	✓
Execute the code of the new process	6	✗
Select another process for execution	7	✗

Your answer is partially correct.

You have correctly selected 1.

The correct answer is: Change to kernel stack → 2, Jump to a code pointed by IDT → 1, Jump to scheduler code → 4, Set the context of the new process → 6, Save the context of the currently running process → 3, Execute the code of the new process → 7, Select another process for execution → 5

Question **17**

Incorrect

Mark 0.00 out of 1.00

Consider the two programs given below to implement the command (ignore the fact that error checks are not done on return values of functions)

\$ ls . /tmp/asdfksdf >/tmp/ddd 2>&1

Program 1

```
int main(int argc, char *argv[]) {
    int fd, n, i;
    char buf[128];

    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    close(1);
    dup(fd);
    close(2);
    dup(fd);
    execl("/bin/ls", "/bin/ls", ".", "/tmp/asldjfaldfs", NULL);
}
```

Program 2

```
int main(int argc, char *argv[]) {
    int fd, n, i;
    char buf[128];

    close(1);
    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    close(2);
    fd = open("/tmp/ddd", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    execl("/bin/ls", "/bin/ls", ".", "/tmp/asldjfaldfs", NULL);
}
```

Select all the correct statements about the programs

Select one or more:

- ☒ a. Both programs are correct ✗
- ☒ b. Program 2 makes sure that there is one file offset used for '2' and '1' ✗
- ☒ c. Only Program 2 is correct ✗
- ☒ d. Program 2 does 1>&2 ✗
- ☒ e. Program 2 ensures 2>&1 and does not ensure > /tmp/ddd ✗
- ☒ f. Program 1 makes sure that there is one file offset used for '2' and '1' ✓
- ☒ g. Program 1 is correct for > /tmp/ddd but not for 2>&1 ✗
- ☒ h. Program 1 does 1>&2 ✗
- ☒ i. Both program 1 and 2 are incorrect ✗
- ☒ j. Program 2 is correct for > /tmp/ddd but not for 2>&1 ✗
- ☒ k. Only Program 1 is correct ✓
- ☒ l. Program 1 ensures 2>&1 and does not ensure > /tmp/ddd ✗

Your answer is incorrect.

The correct answers are: Only Program 1 is correct, Program 1 makes sure that there is one file offset used for '2' and '1'



Question **18**

Correct

Mark 0.25 out of 0.25

Select the option which best describes what the CPU does during it's powered ON lifetime

- ☐ a. Ask the user what is to be done, and execute that task
- ☐ b. Ask the OS what is to be done, and execute that task
- ☐ c. Fetch instructions specified by location given by PC, Decode and Execute it, during execution increment PC or change PC as per the instruction itself, Ask the User or the OS what is to be done next, repeat
- ☒ d. Fetch instructions specified by location given by PC, Decode and Execute it, during execution increment PC or change PC as per the instruction itself, repeat ✓
- ☐ e. Fetch instruction specified by OS, Decode and execute it, repeat
- ☐ f. Fetch instructions specified by location given by PC, Decode and Execute it, during execution increment PC or change PC as per the instruction itself, Ask OS what is to be done next, repeat

The correct answer is: Fetch instructions specified by location given by PC, Decode and Execute it, during execution increment PC or change PC as per the instruction itself, repeat

Question 19

Partially correct

Mark 0.86 out of 1.00

Consider the following code and MAP the file to which each fd points at the end of the code.

```
int main(int argc, char *argv[]) {  
    int fd1, fd2 = 1, fd3 = 1, fd4 = 1;  
  
    fd1 = open("/tmp/1", O_WRONLY | O_CREAT, S_IRUSR|S_IWUSR);  
    fd2 = open("/tmp/2", O_RDONLY);  
    fd3 = open("/tmp/3", O_WRONLY | O_CREAT, S_IRUSR|S_IWUSR);  
    close(0);  
    close(1);  
    dup(fd2);  
    dup(fd3);  
    close(fd3);  
    dup2(fd2, fd4);  
    printf("%d %d %d %d\n", fd1, fd2, fd3, fd4);  
    return 0;  
}
```

1	<input type="text" value="closed"/>	✖
fd4	<input type="text" value="/tmp/2"/>	✔
fd2	<input type="text" value="/tmp/2"/>	✔
fd1	<input type="text" value="/tmp/1"/>	✔
2	<input type="text" value="stderr"/>	✔
0	<input type="text" value="/tmp/2"/>	✔
fd3	<input type="text" value="closed"/>	✔

Your answer is partially correct.

You have correctly selected 6.

The correct answer is: 1 → /tmp/3, fd4 → /tmp/2, fd2 → /tmp/2, fd1 → /tmp/1, 2 → stderr, 0 → /tmp/2, fd3 → closed

Question **20**

Incorrect

Mark 0.00 out of 2.00

Following code claims to implement the command

```
/bin/ls -l | /usr/bin/head -3 | /usr/bin/tail -1
```

Fill in the blanks to make the code work.

Note: Do not include space in writing any option. `x[1][2]` should be written without any space, and so is the case with `[1]` or `[2]`. Pay attention to exact syntax and do not write any extra character like `'` or `=` etc.

```
int main(int argc, char *argv[]) {
```

```
    int pid1, pid2;
```

```
    int pfd[
```

✖ `][2];`

```
    pipe(
```

✖ `);`

```
    pid1 =
```

✖ `;`

```
    if(pid1 != 0) {
```

```
        close(pfd[0]
```

✖ `);`

```
        close(
```

✖ `);`

```
        dup(
```

✖ `);`

```
        execl("/bin/ls", "/bin/ls", "
```

✖ `", NULL);`

```
    }
```

```
    pipe(
```

✖ `);`

✖ `= fork();`

```
    if(pid2 == 0) {
```

```
        close(
```

✖ `;`

```
        close(0);
```

```
        dup(
```

✖ `);`

```
        close(pfd[1]
```

```
✘ );  
    close(  
          
    );  
    dup(  
          
    );  
    execl("/usr/bin/head", "/usr/bin/head", "  
          
    ", NULL);  
    } else {  
        close(pfd  
              
        );  
        close(  
              
        );  
        dup(  
              
        );  
        close(pfd  
              
        );  
        execl("/usr/bin/tail", "/usr/bin/tail", "  
              
        ", NULL);  
    }  
}
```

Question **21**

Partially correct

Mark 0.11 out of 1.00

Select all the correct statements about calling convention on x86 32-bit.

- ☒ a. Return address is one location above the ebp ✓
- ☒ b. Parameters may be passed in registers or on stack ✓
- ☒ c. Space for local variables is allocated by subtracting the stack pointer inside the code of the called function ✓
- ☒ d. The ebp pointers saved on the stack constitute a chain of activation records ✓
- ☒ e. The two lines in the beginning of each function, "push %ebp; mov %esp, %ebp", create space for local variables ✗
- ☒ f. Parameters may be passed in registers or on stack ✓
- ☒ g. The return value is either stored on the stack or returned in the eax register ✗
- ☐ h. Parameters are pushed on the stack in left-right order
- ☐ i. during execution of a function, ebp is pointing to the old ebp
- ☒ j. Space for local variables is allocated by subtracting the stack pointer inside the code of the caller function ✗
- ☒ k. Compiler may allocate more memory on stack than needed ✓

Your answer is partially correct.

You have selected too many options.

The correct answers are: Compiler may allocate more memory on stack than needed, Parameters may be passed in registers or on stack, Parameters may be passed in registers or on stack, Return address is one location above the ebp, during execution of a function, ebp is pointing to the old ebp, Space for local variables is allocated by subtracting the stack pointer inside the code of the called function, The ebp pointers saved on the stack constitute a chain of activation records

Question **22**

Correct

Mark 1.00 out of 1.00

Match the program with its output (ignore newlines in the output. Just focus on the count of the number of 'hi')

- | | | |
|---|-------|---|
| <code>main() { int i = fork(); if(i == 0) execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }</code> | hi | ✓ |
| <code>main() { fork(); execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }</code> | hi hi | ✓ |
| <code>main() { int i = NULL; fork(); printf("hi\n"); }</code> | hi hi | ✓ |
| <code>main() { execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }</code> | hi | ✓ |

Your answer is correct.

The correct answer is: `main() { int i = fork(); if(i == 0) execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }` → hi, `main() { fork(); execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }` → hi hi, `main() { int i = NULL; fork(); printf("hi\n"); }` → hi hi, `main() { execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }` → hi

Question **23**

Incorrect

Mark 0.00 out of 0.50

Some part of the bootloader of xv6 is written in assembly while some part is written in C. Why is that so?

Select all the appropriate choices

- ☒ a. The code in assembly is required for transition to protected mode, from real mode; but calling convention was applicable all the time ✗
- ☒ b. The setting up of the most essential memory management infrastructure needs assembly code ✓
- ☒ c. The code for reading ELF file can not be written in assembly ✗
- ☒ d. The code in assembly is required for transition to protected mode, from real mode; after that calling convention applies, hence code can be written in C ✓

Your answer is incorrect.

The correct answers are: The code in assembly is required for transition to protected mode, from real mode; after that calling convention applies, hence code can be written in C, The setting up of the most essential memory management infrastructure needs assembly code

Question **24**

Incorrect

Mark 0.00 out of 0.50

xv6.img: bootblock kernel

```
dd if=/dev/zero of=xv6.img count=10000
dd if=bootblock of=xv6.img conv=notrunc
dd if=kernel of=xv6.img seek=1 conv=notrunc
```

Consider above lines from the Makefile. Which of the following is incorrect?

- ☒ a. The size of the kernel file is nearly 5 MB ✓
- ☒ b. The kernel is located at block-1 of the xv6.img ✗
- ☒ c. The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies 10,000 blocks on the disk. ✗
- ☐ d. The size of xv6.img is exactly = (size of bootblock) + (size of kernel)
- ☒ e. The bootblock is located on block-0 of the xv6.img ✗
- ☒ f. The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies upto 10,000 blocks on the disk. ✓
- ☒ g. The bootblock may be 512 bytes or less (looking at the Makefile instruction) ✗
- ☒ h. The xv6.img is the virtual disk that is created by combining the bootblock and the kernel file. ✗
- ☒ i. The size of the xv6.img is nearly 5 MB ✗
- ☒ j. xv6.img is the virtual processor used by the qemu emulator ✓
- ☒ k. Blocks in xv6.img after kernel may be all zeroes. ✗

Your answer is incorrect.

The correct answers are: xv6.img is the virtual processor used by the qemu emulator, The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies upto 10,000 blocks on the disk., The size of the kernel file is nearly 5 MB, The size of xv6.img is exactly = (size of bootblock) + (size of kernel)

Question **25**

Incorrect

Mark 0.00 out of 1.00

Select the sequence of events that are NOT possible, assuming a non-interruptible kernel code

Select one or more:

- ☐ a. P1 running
P1 makes system call
timer interrupt
Scheduler
P2 running
timer interrupt
Scheduler
P1 running
P1's system call return
- ☒ b. P1 running
P1 makes system call and blocks
Scheduler
P2 running
P2 makes system call and blocks
Scheduler
P1 running again
- ☐ c. P1 running
P1 makes system call
system call returns
P1 running
timer interrupt
Scheduler running
P2 running
- ☒ d. P1 running
P1 makes system call and blocks
Scheduler
P2 running
P2 makes system call and blocks
Scheduler
P3 running
Hardware interrupt
Interrupt unblocks P1
Interrupt returns
P3 running
Timer interrupt
Scheduler
P1 running
- ☐ e.
P1 running
P1 makes system call
Scheduler
P2 running
P2 makes system call and blocks
Scheduler
P1 running again
- ☒ f. P1 running
keyboard hardware interrupt
keyboard interrupt handler running
interrupt handler returns
P1 running
P1 makes system call
system call returns



P1 running
timer interrupt
scheduler
P2 running

Your answer is incorrect.

The correct answers are: P1 running

P1 makes system call and blocks

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again, P1 running

P1 makes system call

timer interrupt

Scheduler

P2 running

timer interrupt

Scheduler

P1 running

P1's system call return,

P1 running

P1 makes system call

Scheduler

P2 running

P2 makes system call and blocks

Scheduler

P1 running again

Question **26**

Correct

Mark 0.25 out of 0.25

Which of the following are the files related to bootloader in xv6?

- ☐ a. bootasm.s and entry.S
- ☒ b. bootasm.S and bootmain.c
- ☐ c. bootasm.S, bootmain.c and bootblock.c
- ☐ d. bootmain.c and bootblock.S



Your answer is correct.

The correct answer is: bootasm.S and bootmain.c

Question **27**

Correct

Mark 0.25 out of 0.25

Match the following parts of a C program to the layout of the process in memory

Instructions	Text section	✓
Local Variables	Stack Section	✓
Dynamically allocated memory	Heap Section	✓
Global and static data	Data section	✓

Your answer is correct.

The correct answer is:

Instructions → Text section, Local Variables → Stack Section,
Dynamically allocated memory → Heap Section,
Global and static data → Data section

Question **28**

Incorrect

Mark 0.00 out of 0.50

What will this program do?

```
int main() {  
    fork();  
    execl("/bin/ls", "/bin/ls", NULL);  
    printf("hello");  
}
```

- ☐ a. one process will run ls, another will print hello
- ☒ b. run ls once
- ☐ c. run ls twice
- ☐ d. run ls twice and print hello twice
- ☐ e. run ls twice and print hello twice, but output will appear in some random order

✗

Your answer is incorrect.

The correct answer is: run ls twice

Question **29**

Correct

Mark 0.25 out of 0.25

What is the OS Kernel?

- ☒ a. The code that controls hardware, abstracts access to hardware resources using system calls, creates an environment for processes to be created and run ✔ correct
- ☐ b. The set of tools like compiler, linker, loader, terminal, shell, etc.
- ☐ c. Only the system programs like compiler, linker, loader, etc.
- ☐ d. Everything that I see on my screen

The correct answer is: The code that controls hardware, abstracts access to hardware resources using system calls, creates an environment for processes to be created and run

Question **30**

Correct

Mark 0.50 out of 0.50

Which of the following is/are not saved during context switch?

- ☐ a. Program Counter
- ☐ b. General Purpose Registers
- ☒ c. Bus ✔
- ☐ d. Stack Pointer
- ☐ e. MMU related registers/information
- ☒ f. Cache ✔
- ☒ g. TLB ✔

Your answer is correct.

The correct answers are: TLB, Cache, Bus

Question 31

Partially correct

Mark 0.10 out of 0.25

Select the order in which the various stages of a compiler execute.

Linking	3	✗
Syntactical Analysis	2	✓
Pre-processing	1	✓
Intermediate code generation	does not exist	✗
Loading	4	✗

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: Linking → 4, Syntactical Analysis → 2, Pre-processing → 1, Intermediate code generation → 3, Loading → does not exist

Question 32

Partially correct

Mark 0.08 out of 0.50

Order the sequence of events, in scheduling process P1 after process P0

context of P0 is saved in P0's PCB	2	✗
context of P1 is loaded from P1's PCB	3	✗
Process P1 is running	5	✗
timer interrupt occurs	6	✗
Process P0 is running	1	✓
Control is passed to P1	4	✗

Your answer is partially correct.

You have correctly selected 1.

The correct answer is: context of P0 is saved in P0's PCB → 3, context of P1 is loaded from P1's PCB → 4, Process P1 is running → 6, timer interrupt occurs → 2, Process P0 is running → 1, Control is passed to P1 → 5

Question **33**

Not answered

Marked out of 1.00

Select the correct statements about interrupt handling in xv6 code

- ☐ a. On any interrupt/syscall/exception the control first jumps in vectors.S
- ☐ b. The trapframe pointer in struct proc, points to a location on user stack
- ☐ c. Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt
- ☐ d. xv6 uses the 64th entry in IDT for system calls
- ☐ e. The CS and EIP are changed only after pushing user code's SS,ESP on stack
- ☐ f. The trapframe pointer in struct proc, points to a location on kernel stack
- ☐ g. The function trap() is the called only in case of hardware interrupt
- ☐ h. The CS and EIP are changed only immediately on a hardware interrupt
- ☐ i. All the 256 entries in the IDT are filled
- ☐ j. On any interrupt/syscall/exception the control first jumps in trapasm.S
- ☐ k. The function trap() is the called irrespective of hardware interrupt/system-call/exception
- ☐ l. xv6 uses the 0x64th entry in IDT for system calls
- ☐ m. Before going to alltraps, the kernel stack contains upto 5 entries.

Your answer is incorrect.

The correct answers are: All the 256 entries in the IDT are filled, Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt, xv6 uses the 64th entry in IDT for system calls, On any interrupt/syscall/exception the control first jumps in vectors.S, Before going to alltraps, the kernel stack contains upto 5 entries., The trapframe pointer in struct proc, points to a location on kernel stack, The function trap() is the called irrespective of hardware interrupt/system-call/exception, The CS and EIP are changed only after pushing user code's SS,ESP on stack

◀ (Assignment) [Change free list management in xv6](#)

Jump to...

[Dashboard](#) / [My courses](#) / [Computer Engineering & IT](#) / [CEIT-Even-sem-21-22](#) / [OS-even-sem-21-22](#) / [21 March - 27 March](#) / [Quiz-2](#)

Started on Tuesday, 22 March 2022, 1:59:34 PM

State Finished

Completed on Tuesday, 22 March 2022, 4:14:53 PM

Time taken 2 hours 15 mins

Grade 24.57 out of 40.00 (61%)

Question 1

Partially correct

Mark 0.10 out of 1.00

Select all the correct statements w.r.t user and kernel threads

Select one or more:

- ☒ a. many-one model can be implemented even if there are no kernel threads ✓
- ☐ b. many-one model gives no speedup on multicore processors
- ☒ c. all three models, that is many-one, one-one, many-many , require a user level thread library ✓
- ☐ d. A process blocks in many-one model even if a single thread makes a blocking system call
- ☒ e. A process may not block in many-one model, if a thread makes a blocking system call ✗
- ☒ f. one-one model increases kernel's scheduling load ✓
- ☐ g. one-one model can be implemented even if there are no kernel threads

Your answer is partially correct.

You have correctly selected 3.

The correct answers are: many-one model can be implemented even if there are no kernel threads, all three models, that is many-one, one-one, many-many , require a user level thread library, one-one model increases kernel's scheduling load, many-one model gives no speedup on multicore processors, A process blocks in many-one model even if a single thread makes a blocking system call

Question 2

Partially correct

Mark 0.50 out of 1.00

Select all correct statements w.r.t. Major and Minor page faults on Linux

- ☐ a. Thrashing is possible only due to major page faults
- ☐ b. Minor page fault may occur because the page was freed, but still tagged and available in the free page list
- ☒ c. Minor page fault may occur because of a page fault during fork(), on code of an already running process ✓
- ☒ d. Major page faults are likely to occur in more numbers at the beginning of the process ✓
- ☒ e. Minor page fault may occur because the page was a shared memory page ✓
- ☐ f. Minor page faults are an improvement of the page buffering techniques

The correct answers are: Minor page fault may occur because the page was a shared memory page, Minor page fault may occur because of a page fault during fork(), on code of an already running process, Minor page fault may occur because the page was freed, but still tagged and available in the free page list, Major page faults are likely to occur in more numbers at the beginning of the process, Thrashing is possible only due to major page faults, Minor page faults are an improvement of the page buffering techniques

Question 3

Correct

Mark 2.00 out of 2.00

W.r.t. Memory management in xv6,

xv6 uses physical memory upto 224 MB only Mark statements True or False

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	The stack allocated in entry.S is used as stack for scheduler's context for first processor	✓
<input type="radio"/>	<input checked="" type="radio"/>	The switchkvm() call in scheduler() is invoked after control comes to it from swtch() scheduler(), thus demanding execution in new process's context	✓
<input checked="" type="radio"/>	<input type="radio"/>	The switchkvm() call in scheduler() changes CR3 to use page directory kpgdir	✓
<input checked="" type="radio"/>	<input type="radio"/>	The kernel code and data take up less than 2 MB space	✓
<input checked="" type="radio"/>	<input type="radio"/>	The switchkvm() call in scheduler() is invoked after control comes to it from sched(), thus demanding execution in kernel's context	✓
<input checked="" type="radio"/>	<input type="radio"/>	xv6 uses physical memory upto 224 MB only	✓
<input checked="" type="radio"/>	<input type="radio"/>	The free page-frame are created out of nearly 222 MB	✓
<input type="radio"/>	<input checked="" type="radio"/>	The switchkvm() call in scheduler() changes CR3 to use page directory of new process	✓
<input checked="" type="radio"/>	<input type="radio"/>	PHYSTOP can be increased to some extent, simply by editing memlayout.h	✓
<input checked="" type="radio"/>	<input type="radio"/>	The process's address space gets mapped on frames, obtained from ~2MB:224MB range	✓
<input type="radio"/>	<input checked="" type="radio"/>	The kernel's page table given by kpgdir variable is used as stack for scheduler's context	✓

The stack allocated in entry.S is used as stack for scheduler's context for first processor: True

The switchkvm() call in scheduler() is invoked after control comes to it from swtch() scheduler(), thus demanding execution in new process's context: False

The switchkvm() call in scheduler() changes CR3 to use page directory kpgdir: True

The kernel code and data take up less than 2 MB space: True

The switchkvm() call in scheduler() is invoked after control comes to it from sched(), thus demanding execution in kernel's context: True

xv6 uses physical memory upto 224 MB only: True

The free page-frame are created out of nearly 222 MB: True

The switchkvm() call in scheduler() changes CR3 to use page directory of new process: False

PHYSTOP can be increased to some extent, simply by editing memlayout.h: True

The process's address space gets mapped on frames, obtained from ~2MB:224MB range: True

The kernel's page table given by kpgdir variable is used as stack for scheduler's context: False

Question 4

Correct

Mark 1.00 out of 1.00

Given that a kernel has 1000 KB of total memory, and holes of sizes (in that order) 300 KB, 200 KB, 100 KB, 250 KB. For each of the requests on the left side, match it with the chunk chosen using the specified algorithm.

Consider each request as first request.

50 KB, worst fit	300 KB	✓
100 KB, worst fit	300 KB	✓
200 KB, first fit	300 KB	✓
150 KB, best fit	200 KB	✓
220 KB, best fit	250 KB	✓
150 KB, first fit	300 KB	✓

The correct answer is: 50 KB, worst fit → 300 KB, 100 KB, worst fit → 300 KB, 200 KB, first fit → 300 KB, 150 KB, best fit → 200 KB, 220 KB, best fit → 250 KB, 150 KB, first fit → 300 KB

Question 5

Partially correct

Mark 0.60 out of 1.00

Choice of the global or local replacement strategy is a subjective choice for kernel programmers. There are advantages and disadvantages on either side. Out of the following statements, that advocate either global or local replacement strategy, select those statements that have a logically CONSISTENT argument. (That is any statement that is logically correct about either global or local replacement)

Consistent	Inconsistent		
<input checked="" type="radio"/>	<input type="radio"/>	Local replacement can lead to under-utilisation of memory, because a process may not use all the pages allocated to it all the time.	✓
<input type="radio"/>	<input checked="" type="radio"/>	Global replacement may give highly variable per process completion time because number of page faults become un-predictable.	✗
<input checked="" type="radio"/>	<input type="radio"/>	Global replacement can be preferred when greater throughput (number of processes completing per unit time) is a concern, because each process tries to complete at the expense of others, thus leading to overall more processes completing (unless thrashing occurs).	✓
<input checked="" type="radio"/>	<input type="radio"/>	Local replacement results in more predictable per-process completion time because number of page faults can be better predicted.	✓
<input type="radio"/>	<input checked="" type="radio"/>	Local replacement can be preferred when avoiding thrashing is a major concern because with local replacement and minimum number of frames allocated, a process is always able to progress and cascading inter-process page faults are avoided.	✗

Local replacement can lead to under-utilisation of memory, because a process may not use all the pages allocated to it all the time.: Consistent

Global replacement may give highly variable per process completion time because number of page faults become un-predictable.: Consistent

Global replacement can be preferred when greater throughput (number of processes completing per unit time) is a concern, because each process tries to complete at the expense of others, thus leading to overall more processes completing (unless thrashing occurs).: Consistent

Local replacement results in more predictable per-process completion time because number of page faults can be better predicted.: Consistent

Local replacement can be preferred when avoiding thrashing is a major concern because with local replacement and minimum number of frames allocated, a process is always able to progress and cascading inter-process page faults are avoided.: Consistent

Question 6

Correct

Mark 1.00 out of 1.00

Map the functionality/use with function/variable in xv6 code.

Setup kernel part of a page table, and switch to that page table

kvmmalloc()



Setup kernel part of a page table, mapping kernel code, data, read-only data, I/O space, devices

setupkvm()



return a free page, if available; 0, otherwise

kalloc()



Create page table entries for a given range of virtual and physical addresses; including page directory entries if needed

mappages()



Return address of page table entry in a given page directory, for a given virtual address; creates page table if necessary

walkpgdir()



Array listing the kernel memory mappings, to be used by setupkvm()

kmap[]



Your answer is correct.

The correct answer is: Setup kernel part of a page table, and switch to that page table → kvmmalloc(), Setup kernel part of a page table, mapping kernel code, data, read-only data, I/O space, devices → setupkvm(), return a free page, if available; 0, otherwise → kalloc(), Create page table entries for a given range of virtual and physical addresses; including page directory entries if needed → mappages(), Return address of page table entry in a given page directory, for a given virtual address; creates page table if necessary → walkpgdir(), Array listing the kernel memory mappings, to be used by setupkvm() → kmap[]

Question 7

Partially correct

Mark 0.20 out of 1.00

Mark statements True/False w.r.t. change of states of a process. Note that a statement is true only if the claim and argument both are true.

Reference: The process state diagram (and your understanding of how kernel code works). Note - the diagram does not show zombie state!




Figure 3.2 Diagram of process state.

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	Every forked process has to go through ZOMBIE state, at least for a small duration.	✓
<input type="radio"/>	<input checked="" type="radio"/>	A process only in RUNNING state can become TERMINATED because scheduler moves it to ZOMBIE state first	✗
<input checked="" type="radio"/>	<input type="radio"/>	A process in WAITING state can not become RUNNING because the event it's waiting for has not occurred and it has not been moved to ready queue yet	✗
<input checked="" type="radio"/>	<input type="radio"/>	Only a process in READY state is considered by scheduler	✗
<input type="radio"/>	<input checked="" type="radio"/>	A process in READY state can not go to WAITING state because the resource on which it will WAIT will not be in use when process is in READY state.	✗

Every forked process has to go through ZOMBIE state, at least for a small duration.: True

A process only in RUNNING state can become TERMINATED because scheduler moves it to ZOMBIE state first: False

A process in WAITING state can not become RUNNING because the event it's waiting for has not occurred and it has not been moved to ready queue yet: True

Only a process in READY state is considered by scheduler: True

A process in READY state can not go to WAITING state because the resource on which it will WAIT will not be in use when process is in READY state.: False

Question 8

Correct

Mark 2.00 out of 2.00

Consider the reference string

6 4 2 0 1 2 6 9 2 0 5

If the number of page frames is 3, then total number of page faults (including initial), using FIFO replacement is:

Answer:

10



#6# 6,4# 6,4,2 #0,4,2# 0,1,2 #0,1,6 #9,1,6# 9,2,6# 9,2,0 #5,2,0

The correct answer is: 10

Question 9

Incorrect

Mark 0.00 out of 1.00

Select all the correct statements about linking and loading.

Select one or more:

- ☐ a. Continuous memory management schemes can support static linking and dynamic loading. (may be inefficiently)
- ☐ b. Dynamic linking essentially results in relocatable code.
- ☐ c. Continuous memory management schemes can support static linking and static loading. (may be inefficiently)
- ☐ d. Loader is last stage of the linker program
- ☐ e. Dynamic linking and loading is not possible without demand paging or demand segmentation.
- ☐ f. Static linking leads to non-relocatable code
- ☐ g. Continuous memory management schemes can support dynamic linking and dynamic loading.
- ☒ h. Dynamic linking is possible with continuous memory management, but variable sized partitions only. ✗
- ☒ i. Loader is part of the operating system ✓

Your answer is incorrect.

The correct answers are: Continuous memory management schemes can support static linking and static loading. (may be inefficiently), Continuous memory management schemes can support static linking and dynamic loading. (may be inefficiently), Dynamic linking essentially results in relocatable code., Loader is part of the operating system, Dynamic linking and loading is not possible without demand paging or demand segmentation.

Question 10

Correct

Mark 1.00 out of 1.00

Consider a computer system with a 32-bit logical address and 4- KB page size. The system supports up to 512 MB of physical memory. How many entries are there in each of the following?

Write answer as a decimal number.

A conventional, single-level page table:



An inverted page table:



Question 11

Incorrect

Mark 0.00 out of 1.00

W.r.t. xv6 code, match the state of a process with a code that sets the state

RUNNING	<input type="text" value="Choose..."/>	
ZOMBIE	<input type="text" value="Choose..."/>	
EMBRYO	<input type="text" value="Choose..."/>	
SLEEPING	<input type="text" value="Choose..."/>	
UNUSED	<input type="text" value="Choose..."/>	
RUNNABLE	<input type="text" value="scheduler()"/>	✗

The correct answer is: RUNNING → scheduler(), ZOMBIE → exit(), called by process itself, EMBRYO → fork()->allocproc() before setting up the UVM, SLEEPING → sleep(), called by any process blocking itself, UNUSED → wait(), called by parent process, RUNNABLE → wakeup(), called by an interrupt handler

Question **12**

Not answered

Marked out of 1.00

Select the correct statements about interrupt handling in xv6 code

- ☐ a. The trapframe pointer in struct proc, points to a location on user stack
- ☐ b. The CS and EIP are changed only immediately on a hardware interrupt
- ☐ c. The CS and EIP are changed only after pushing user code's SS,ESP on stack
- ☐ d. The trapframe pointer in struct proc, points to a location on kernel stack
- ☐ e. On any interrupt/syscall/exception the control first jumps in vectors.S
- ☐ f. The function trap() is the called irrespective of hardware interrupt/system-call/exception
- ☐ g. All the 256 entries in the IDT are filled
- ☐ h. The function trap() is the called only in case of hardware interrupt
- ☐ i. xv6 uses the 64th entry in IDT for system calls
- ☐ j. xv6 uses the 0x64th entry in IDT for system calls
- ☐ k. Before going to alltraps, the kernel stack contains upto 5 entries.
- ☐ l. Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt
- ☐ m. On any interrupt/syscall/exception the control first jumps in trapasm.S

Your answer is incorrect.

The correct answers are: All the 256 entries in the IDT are filled, Each entry in IDT essentially gives the values of CS and EIP to be used in handling that interrupt, xv6 uses the 64th entry in IDT for system calls, On any interrupt/syscall/exception the control first jumps in vectors.S, Before going to alltraps, the kernel stack contains upto 5 entries., The trapframe pointer in struct proc, points to a location on kernel stack, The function trap() is the called irrespective of hardware interrupt/system-call/exception, The CS and EIP are changed only after pushing user code's SS,ESP on stack

Question 13

Correct

Mark 1.00 out of 1.00

The complete range of virtual addresses (after main() in main.c is over), from which the free pages used by kalloc() and kfree() is derived, are:

- ☐ a. end, (4MB + PHYSTOP)
- ☐ b. P2V(end), PHYSTOP
- ☐ c. end, P2V(4MB + PHYSTOP)
- ☐ d. end, PHYSTOP
- ☐ e. end, 4MB
- ☒ f. end, P2V(PHYSTOP)
- ☐ g. P2V(end), P2V(PHYSTOP)



Your answer is correct.

The correct answer is: end, P2V(PHYSTOP)

Question 14

Correct

Mark 1.00 out of 1.00

Select all the correct statements about MMU and its functionality (on a non-demand paged system)

Select one or more:

- ☒ a. Illegal memory access is detected in hardware by MMU and a trap is raised
- ☐ b. Illegal memory access is detected by operating system
- ☐ c. MMU is a separate chip outside the processor
- ☐ d. The operating system interacts with MMU for every single address translation
- ☒ e. Logical to physical address translations in MMU are done in hardware, automatically
- ☐ f. Logical to physical address translations in MMU are done with specific machine instructions
- ☒ g. MMU is inside the processor
- ☒ h. The Operating system sets up relevant CPU registers to enable proper MMU translations



Your answer is correct.

The correct answers are: MMU is inside the processor, Logical to physical address translations in MMU are done in hardware, automatically, The Operating system sets up relevant CPU registers to enable proper MMU translations, Illegal memory access is detected in hardware by MMU and a trap is raised

Question **15**

Incorrect

Mark 0.00 out of 2.00

Order the following events, in the creation of init() process in xv6:

1. ☒ initcode is selected by scheduler for execution
2. ☒ kernel stack is allocated for initcode process
3. ☒ values are set in the trapframe of initcode
4. ☒ sys_exec runs
5. ☒ initcode process is set to be runnable
6. ☒ code is set to start in forkret() when process gets scheduled
7. ☒ Arguments on setup on process stack for /init
8. ☒ trapframe and context pointers are set to proper location
9. ☒ trap() runs
10. ☒ userinit() is called
11. ☒ the header of "/init" ELF file is ready by kernel
12. ☒ empty struct proc is obtained for initcode
13. ☒ Stack is allocated for "/init" process
14. ☒ function pointer from syscalls[] array is invoked
15. ☒ memory mappings are created for "/init" process
16. ☒ page table mappings of 'initcode' are replaced by mappings of 'init'
17. ☒ initcode calls exec system call
18. ☒ initcode process runs
19. ☒ name of process "/init" is copied in struct proc

Your answer is incorrect.

Grading type: Relative to the next item (including last)

Grade details: 0 / 19 = 0%

Here are the scores for each item in this response:

1. 0 / 1 = 0%
2. 0 / 1 = 0%
3. 0 / 1 = 0%
4. 0 / 1 = 0%
5. 0 / 1 = 0%
6. 0 / 1 = 0%
7. 0 / 1 = 0%
8. 0 / 1 = 0%
9. 0 / 1 = 0%
10. 0 / 1 = 0%
11. 0 / 1 = 0%
12. 0 / 1 = 0%
13. 0 / 1 = 0%
14. 0 / 1 = 0%
15. 0 / 1 = 0%

- 16. 0 / 1 = 0%
- 17. 0 / 1 = 0%
- 18. 0 / 1 = 0%
- 19. 0 / 1 = 0%

The correct order for these items is as follows:

1. userinit() is called
2. empty struct proc is obtained for initcode
3. kernel stack is allocated for initcode process
4. trapframe and context pointers are set to proper location
5. code is set to start in forkret() when process gets scheduled
6. kernel memory mappings are created for initcode
7. values are set in the trapframe of initcode
8. initcode process is set to be runnable
9. initcode is selected by scheduler for execution
10. initcode process runs
11. initcode calls exec system call
12. trap() runs
13. function pointer from syscalls[] array is invoked
14. sys_exec runs
15. the header of "/init" ELF file is ready by kernel
16. memory mappings are created for "/init" process
17. Stack is allocated for "/init" process
18. Arguments on setup on process stack for /init
19. name of process "/init" is copied in struct proc
20. page table mappings of 'initcode' are replaced by mappings of 'init'

Question **16**

Partially correct

Mark 0.56 out of 1.00

Mark the statements as True or False, w.r.t. mmap()

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	mmap() can be implemented on both demand paged and non-demand paged systems.	✗
<input type="radio"/>	<input checked="" type="radio"/>	MAP_FIXED guarantees that the mapping is always done at the specified address	✗
<input checked="" type="radio"/>	<input type="radio"/>	MAP_PRIVATE leads to a mapping that is copy-on-write	✓
<input checked="" type="radio"/>	<input type="radio"/>	on failure mmap() returns (void *)-1	✓
<input type="radio"/>	<input checked="" type="radio"/>	MAP_SHARED leads to a mapping that is copy-on-write	✓
<input type="radio"/>	<input checked="" type="radio"/>	mmap() results in changes to buffer-cache of the kernel.	✗
<input type="radio"/>	<input checked="" type="radio"/>	on failure mmap() returns NULL	✓
<input checked="" type="radio"/>	<input type="radio"/>	mmap() results in changes to page table of a process.	✗
<input checked="" type="radio"/>	<input type="radio"/>	mmap() is a system call	✓

mmap() can be implemented on both demand paged and non-demand paged systems.: True

MAP_FIXED guarantees that the mapping is always done at the specified address: False

MAP_PRIVATE leads to a mapping that is copy-on-write: True

on failure mmap() returns (void *)-1: True

MAP_SHARED leads to a mapping that is copy-on-write: False

mmap() results in changes to buffer-cache of the kernel.: False

on failure mmap() returns NULL: False

mmap() results in changes to page table of a process.: True

mmap() is a system call: True

Question **17**

Incorrect

Mark 0.00 out of 1.00

If one thread opens a file with read privileges then

Select one:

- ☐ a. other threads in the same process can also read from that file
- ☐ b. none of these
- ☐ c. any other thread cannot read from that file
- ☒ d. other threads in the another process can also read from that file



Your answer is incorrect.

The correct answer is: other threads in the same process can also read from that file

Question 18

Partially correct

Mark 0.60 out of 1.00

Mark the statements about named and un-named pipes as True or False

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	Named pipe exists as a file	✓
<input checked="" type="radio"/>	<input type="radio"/>	Un-named pipes are inherited by a child process from parent.	✓
<input type="radio"/>	<input checked="" type="radio"/>	The buffers for named-pipe are in process-memory while the buffers for the un-named pipe are in kernel memory.	✗
<input checked="" type="radio"/>	<input type="radio"/>	Both types of pipes are an extension of the idea of "message passing".	✓
<input type="radio"/>	<input checked="" type="radio"/>	A named pipe has a name decided by the kernel.	✗
<input checked="" type="radio"/>	<input type="radio"/>	Un-named pipes can be used for communication between only "related" processes, if the common ancestor created it.	✗
<input checked="" type="radio"/>	<input type="radio"/>	Both types of pipes provide FIFO communication.	✓
<input type="radio"/>	<input checked="" type="radio"/>	Named pipes can be used for communication between only "related" processes.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Named pipes can exist beyond the life-time of processes using them.	✓
<input type="radio"/>	<input checked="" type="radio"/>	The pipe() system call can be used to create either a named or un-named pipe.	✗

Named pipe exists as a file: True

Un-named pipes are inherited by a child process from parent.: True

The buffers for named-pipe are in process-memory while the buffers for the un-named pipe are in kernel memory.: False

Both types of pipes are an extension of the idea of "message passing".: True

A named pipe has a name decided by the kernel.: False

Un-named pipes can be used for communication between only "related" processes, if the common ancestor created it.: True

Both types of pipes provide FIFO communication.: True

Named pipes can be used for communication between only "related" processes.: False

Named pipes can exist beyond the life-time of processes using them.: True

The pipe() system call can be used to create either a named or un-named pipe.: False

Question **19**

Partially correct

Mark 0.67 out of 1.00

Select the most common causes of use of IPC by processes

- ☐ a. More modular code
- ☒ b. Breaking up a large task into small tasks and speeding up computation, on multiple core machines ✓
- ☐ c. More security checks
- ☒ d. Sharing of information of common interest ✓
- ☐ e. Get the kernel performance statistics

The correct answers are: Sharing of information of common interest, Breaking up a large task into small tasks and speeding up computation, on multiple core machines, More modular code

Question **20**

Correct

Mark 1.00 out of 1.00

For each function/code-point, select the status of segmentation setup in xv6

after seginit() in main()	gdt setup with 5 entries (0 to 4) on one processor	✓
bootmain()	gdt setup with 3 entries, at start32 symbol of bootasm.S	✓
after startothers() in main()	gdt setup with 5 entries (0 to 4) on all processors	✓
entry.S	gdt setup with 3 entries, at start32 symbol of bootasm.S	✓
kvmalloc() in main()	gdt setup with 3 entries, at start32 symbol of bootasm.S	✓
bootasm.S	gdt setup with 3 entries, at start32 symbol of bootasm.S	✓

Your answer is correct.

The correct answer is: after seginit() in main() → gdt setup with 5 entries (0 to 4) on one processor, bootmain() → gdt setup with 3 entries, at start32 symbol of bootasm.S, after startothers() in main() → gdt setup with 5 entries (0 to 4) on all processors, entry.S → gdt setup with 3 entries, at start32 symbol of bootasm.S, kvmalloc() in main() → gdt setup with 3 entries, at start32 symbol of bootasm.S, bootasm.S → gdt setup with 3 entries, at start32 symbol of bootasm.S

Question **21**

Partially correct

Mark 0.50 out of 1.00

Mark whether the given sequence of events is possible or not-possible. Also, select the reason for your answer.

For each sequence it's a not-possible sequence if some important event is not mentioned in the sequence.

Assume that the kernel code is non-interruptible and uniprocessor system.

Process P1, user code executing

Timer interrupt

Context changes to kernel context

Generic interrupt handler runs


Generic interrupt handler calls Scheduler

Scheduler selects P2 for execution

After scheduler, Process P2 user code executing

This sequence of events is: 

Because



Question **22**

Partially correct

Mark 0.63 out of 1.00

Mark the statements as True or False, w.r.t. passing of arguments to system calls in xv6 code.

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	Integer arguments are stored in eax, ebx, ecx, etc. registers	✗
<input checked="" type="radio"/>	<input type="radio"/>	String arguments are NOT copied in kernel memory, but just pointed to by a kernel memory pointer	✓
<input checked="" type="radio"/>	<input type="radio"/>	The functions like argint(), argstr() make the system call arguments available in the kernel.	✓
<input checked="" type="radio"/>	<input type="radio"/>	String arguments are first copied to trapframe and then from trapframe to kernel's other variables.	✗
<input checked="" type="radio"/>	<input type="radio"/>	The arguments to system call originally reside on process stack.	✓
<input checked="" type="radio"/>	<input type="radio"/>	The arguments to system call are copied to kernel stack in trapasm.S	✗
<input checked="" type="radio"/>	<input type="radio"/>	Integer arguments are copied from user memory to kernel memory using argint()	✓
<input checked="" type="radio"/>	<input type="radio"/>	The arguments are accessed in the kernel code using esp on the trapframe.	✓

Integer arguments are stored in eax, ebx, ecx, etc. registers: False

String arguments are NOT copied in kernel memory, but just pointed to by a kernel memory pointer: True

The functions like argint(), argstr() make the system call arguments available in the kernel.: True

String arguments are first copied to trapframe and then from trapframe to kernel's other variables.: False

The arguments to system call originally reside on process stack.: True

The arguments to system call are copied to kernel stack in trapasm.S: False

Integer arguments are copied from user memory to kernel memory using argint(): True

The arguments are accessed in the kernel code using esp on the trapframe.: True

Question **23**

Not answered

Marked out of 1.00

Given below is a sequence of reference bits on pages before the second chance algorithm runs. Before the algorithm runs, the counter is at the page marked (x). Write the sequence of reference bits after the second chance algorithm has executed once. In the answer write PRECISELY one space BETWEEN each number and do not mention (x).

0 0 1(x) 1 0 1 1

Answer:



The correct answer is: 0 0 0 0 1 1

Question **24**

Correct

Mark 2.00 out of 2.00

For the reference string

3 4 3 5 2

using FIFO replacement policy for pages,

consider the number of page faults for 2, 3 and 4 page frames.

Select the correct statement.

Select one:

- ☐ a. Exhibit Balady's anomaly between 3 and 4 frames
- ☒ b. Do not exhibit Balady's anomaly
- ☐ c. Exhibit Balady's anomaly between 2 and 3 frames



Your answer is correct.

The correct answer is: Do not exhibit Balady's anomaly

Question **25**

Correct

Mark 1.00 out of 1.00

For the reference string

3 4 3 5 2

using LRU replacement policy for pages,

consider the number of page faults for 2, 3 and 4 page frames.

Select the most correct statement.

Select one:

- ☒ a. LRU will never exhibit Balady's anomaly
- ☐ b. Exhibit Balady's anomaly between 2 and 3 frames
- ☐ c. This example does not exhibit Balady's anomaly
- ☐ d. Exhibit Balady's anomaly between 3 and 4 frames



Your answer is correct.

The correct answer is: LRU will never exhibit Balady's anomaly

Question **26**

Partially correct

Mark 0.55 out of 1.00

Select all the correct statements about process states.

Note that in this question you lose marks for every incorrect choice that you make, proportional to actual number of incorrect choices.

- ☐ a. Process state is implemented as a string
- ☒ b. Process state is stored in the PCB ✓
- ☐ c. A process becomes ZOMBIE when another process bites into it's memory
- ☒ d. Process state is stored in the processor ✗
- ☐ e. The scheduler can change state of a process from RUNNABLE to RUNNING and vice-versa
- ☒ f. The scheduler can change state of a process from RUNNABLE to RUNNING ✓
- ☒ g. A process becomes ZOMBIE when it calls exit() ✓
- ☐ h. Process state is changed only by interrupt handlers
- ☐ i. Process state can be implemented as just a number

Your answer is partially correct.

You have correctly selected 3.

The correct answers are: Process state is stored in the PCB, Process state can be implemented as just a number, The scheduler can change state of a process from RUNNABLE to RUNNING, A process becomes ZOMBIE when it calls exit()

Question 27

Partially correct

Mark 0.38 out of 1.00

Consider a demand-paging system with the following time-measured utilizations:

CPU utilization : 20%

Paging disk: 97.7%

Other I/O devices: 5%

For each of the following, indicate whether it will (or is likely to) improve CPU utilization (even if by a small amount). Explain your answers.

- a. Install a faster CPU : Yes ☐ ✖
- b. Install a bigger paging disk. : Yes ☐ ✖
- c. Increase the degree of multiprogramming. : Yes ☐ ✖
- d. Decrease the degree of multiprogramming. : Yes ☐ ✔
- e. Install more main memory.: Yes ☐ ✔
- f. Install a faster hard disk or multiple controllers with multiple hard disks. : Yes ☐ ✔
- g. Add prepaging to the page-fetch algorithms. : ✖
- h. Increase the page size. : ✖

Question 28

Incorrect

Mark 0.00 out of 1.00

Suppose a kernel uses a buddy allocator. The smallest chunk that can be allocated is of size 32 bytes. One bit is used to track each such chunk, where 1 means allocated and 0 means free. The chunk looks like this as of now:

10011010

Now, there is a request for a chunk of 50 bytes.

After this allocation, the bitmap, indicating the status of the buddy allocator will be

Answer: ✖

The correct answer is: 11111010

Question **29**

Partially correct

Mark 0.75 out of 1.00

Select the correct points of comparison between POSIX and System V shared memory.

- ☒ a. POSIX shared memory is newer than System V shared memory ✓
- ☒ b. POSIX shared memory is "thread safe", System V is not ✓
- ☒ c. System V is more prevalent than POSIX even today ✓
- ☐ d. POSIX allows giving name to shared memory, System V does not

The correct answers are: POSIX shared memory is newer than System V shared memory, POSIX shared memory is "thread safe", System V is not, POSIX allows giving name to shared memory, System V does not, System V is more prevalent than POSIX even today

Question 30

Partially correct

Mark 0.67 out of 1.00




Figure 9.8 Paging hardware.

Mark the statements as True or False, w.r.t. the above diagram (note that the diagram does not cover all details of what actually happens!)

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	The combining of f and d is done by MMU	✓
<input type="radio"/>	<input checked="" type="radio"/>	There are total 3 memory references in this diagram	✗
<input checked="" type="radio"/>	<input type="radio"/>	The split of logical address into p and d is done by MMU	✓
<input checked="" type="radio"/>	<input type="radio"/>	The page table is in physical memory and must be continuous	✓
<input type="radio"/>	<input checked="" type="radio"/>	Using the offset d in the physical page-frame is done by MMU	✗
<input checked="" type="radio"/>	<input type="radio"/>	The logical address issued by CPU is the same one generated by compiler	✓

The combining of f and d is done by MMU: True

There are total 3 memory references in this diagram: False

The split of logical address into p and d is done by MMU: True

The page table is in physical memory and must be continuous: True

Using the offset d in the physical page-frame is done by MMU: False

The logical address issued by CPU is the same one generated by compiler: True

Question **31**

Partially correct

Mark 0.50 out of 1.00

Select all the correct statements about signals

Select one or more:

- ☐ a. SIGKILL definitely kills a process because it's code runs in kernel mode of CPU
- ☒ b. Signals are delivered to a process by another process ✗
- ☐ c. The signal handler code runs in kernel mode of CPU
- ☒ d. SIGKILL definitely kills a process because it can't be caught or ignored, and it's default action terminates the process ✓
- ☒ e. The signal handler code runs in user mode of CPU ✓
- ☒ f. A signal handler can be invoked asynchronously or synchronously depending on signal type ✓
- ☐ g. Signal handlers once replaced can't be restored
- ☐ h. Signals are delivered to a process by kernel

Your answer is partially correct.

You have correctly selected 3.

The correct answers are: Signals are delivered to a process by kernel, A signal handler can be invoked asynchronously or synchronously depending on signal type, The signal handler code runs in user mode of CPU, SIGKILL definitely kills a process because it can't be caught or ignored, and it's default action terminates the process

Question **32**

Correct

Mark 1.00 out of 1.00

The data structure used in `kalloc()` and `kfree()` in xv6 is

- ☐ a. Singly linked circular list
- ☒ b. Singly linked NULL terminated list ✓
- ☐ c. Double linked NULL terminated list
- ☐ d. Doubly linked circular list

Your answer is correct.

The correct answer is: Singly linked NULL terminated list

Question **33**

Partially correct

Mark 1.78 out of 2.00

Match the description of a memory management function with the name of the function that provides it, in xv6

Load contents from ELF into existing pages

loaduvm()



Mark the page as in-accessible

clearpteu()



setup the kernel part in the page table

setupkvm()



Switch to kernel page table

switchkvm()



Create a copy of the page table of a process

copyuvm()



Copy the code pages of a process

No such function



Setup and load the user page table for initcode process

inituvm()



Switch to user page table

switchuvm()



Load contents from ELF into pages after allocating the pages first

inituvm()



The correct answer is: Load contents from ELF into existing pages → loaduvm(), Mark the page as in-accessible → clearpteu(), setup the kernel part in the page table → setupkvm(), Switch to kernel page table → switchkvm(), Create a copy of the page table of a process → copyuvm(), Copy the code pages of a process → No such function, Setup and load the user page table for initcode process → inituvm(), Switch to user page table → switchuvm(), Load contents from ELF into pages after allocating the pages first → No such function

Question 34

Partially correct

Mark 0.60 out of 1.00

Mark the statements as True or False, w.r.t. thrashing

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	Thrashing occurs because some process is doing lot of disk I/O.	✗
<input checked="" type="radio"/>	<input type="radio"/>	Processes keep changing their locality of reference, and a high rate of page faults occur when they are changing the locality.	✓
<input type="radio"/>	<input checked="" type="radio"/>	mmap() solves the problem of thrashing.	✓
<input checked="" type="radio"/>	<input type="radio"/>	The working set model is an attempt at approximating the locality of a process.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Thrashing is particular to demand paging systems, and does not apply to pure paging systems.	✓
<input type="radio"/>	<input checked="" type="radio"/>	Processes keep changing their locality of reference, and least number of page faults occur when they are changing the locality.	✗
<input type="radio"/>	<input checked="" type="radio"/>	Thrashing can occur even if entire memory is not in use.	✗
<input checked="" type="radio"/>	<input type="radio"/>	During thrashing the CPU is under-utilised as most time is spent in I/O	✗
<input checked="" type="radio"/>	<input type="radio"/>	Thrashing can be limited if local replacement is used.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Thrashing occurs when the total size of all processes's locality exceeds total memory size.	✓

Thrashing occurs because some process is doing lot of disk I/O.: False

Processes keep changing their locality of reference, and a high rate of page faults occur when they are changing the locality.: True

mmap() solves the problem of thrashing.: False

The working set model is an attempt at approximating the locality of a process.: True

Thrashing is particular to demand paging systems, and does not apply to pure paging systems.: True

Processes keep changing their locality of reference, and least number of page faults occur when they are changing the locality.: False

Thrashing can occur even if entire memory is not in use.: False

During thrashing the CPU is under-utilised as most time is spent in I/O: True

Thrashing can be limited if local replacement is used.: True

Thrashing occurs when the total size of all processes's locality exceeds total memory size.: True

Question **35**

Correct

Mark 1.00 out of 1.00

After virtual memory is implemented

(select T/F for each of the following) One Program's size can be larger than physical memory size

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	Cumulative size of all programs can be larger than physical memory size	✓
<input checked="" type="radio"/>	<input type="radio"/>	Code need not be completely in memory	✓
<input checked="" type="radio"/>	<input type="radio"/>	One Program's size can be larger than physical memory size	✓
<input type="radio"/>	<input checked="" type="radio"/>	Virtual addresses become available to executing process	✓
<input type="radio"/>	<input checked="" type="radio"/>	Virtual access to memory is granted to all processes	✓
<input checked="" type="radio"/>	<input type="radio"/>	Relatively less I/O may be possible during process execution	✓
<input checked="" type="radio"/>	<input type="radio"/>	Logical address space could be larger than physical address space	✓

Cumulative size of all programs can be larger than physical memory size: True

Code need not be completely in memory: True

One Program's size can be larger than physical memory size: True

Virtual addresses become available to executing process: False

Virtual access to memory is granted to all processes: False

Relatively less I/O may be possible during process execution: True

Logical address space could be larger than physical address space: True

[◀ \(Optional Assignment\) lseek system call in xv6](#)

Jump to...

[Feedback on Quiz-2 ►](#)

[Dashboard](#) / [My courses](#) / [Computer Engineering & IT](#) / [CEIT-Even-sem-20-21](#) / [OS-Even-sem-2020-21](#) / 14 March - 20 March
/ [Quiz - 2 \(18 March\)](#)

Started on Thursday, 18 March 2021, 2:46 PM

State Finished

Completed on Thursday, 18 March 2021, 3:50 PM

Time taken 1 hour 4 mins

Grade 10.36 out of 20.00 (52%)

Question 1

Partially correct

Mark 0.57 out of 1.00

Mark True, the actions done as part of code of switch() in switch.S, in xv6

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	Restore new callee saved registers from kernel stack of new context	✓
<input checked="" type="radio"/>	<input type="radio"/>	Save old callee saved registers on kernel stack of old context	✓
<input type="radio"/>	<input checked="" type="radio"/>	Save old callee saved registers on user stack of old context	✓
<input type="radio"/>	<input checked="" type="radio"/>	Switch from old process context to new process context	✗
<input checked="" type="radio"/>	<input type="radio"/>	Switch from one stack (old) to another(new)	✗
<input type="radio"/>	<input checked="" type="radio"/>	Restore new callee saved registers from user stack of new context	✓
<input type="radio"/>	<input checked="" type="radio"/>	Jump to code in new context	✗

Restore new callee saved registers from kernel stack of new context: True

Save old callee saved registers on kernel stack of old context: True

Save old callee saved registers on user stack of old context: False

Switch from old process context to new process context: False

Switch from one stack (old) to another(new): True

Restore new callee saved registers from user stack of new context: False

Jump to code in new context: False

Question 2

Partially correct

Mark 0.17 out of 0.50

For each function/code-point, select the status of segmentation setup in xv6

bootmain()	gdt setup with 3 entries, right from first line of code of bootloader	✗
kvmalloc() in main()	gdt setup with 5 entries (0 to 4) on one processor	✗
after startothers() in main()	gdt setup with 5 entries (0 to 4) on all processors	✓
after seginit() in main()	gdt setup with 5 entries (0 to 4) on all processors	✗
bootasm.S	gdt setup with 3 entries, right from first line of code of bootloader	✗
entry.S	gdt setup with 3 entries, at start32 symbol of bootasm.S	✓

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: bootmain() → gdt setup with 3 entries, at start32 symbol of bootasm.S, kvmalloc() in main() → gdt setup with 3 entries, at start32 symbol of bootasm.S, after startothers() in main() → gdt setup with 5 entries (0 to 4) on all processors, after seginit() in main() → gdt setup with 5 entries (0 to 4) on one processor, bootasm.S → gdt setup with 3 entries, at start32 symbol of bootasm.S, entry.S → gdt setup with 3 entries, at start32 symbol of bootasm.S

Question 3

Partially correct

Mark 0.38 out of 1.00

Compare paging with demand paging and select the correct statements.

Select one or more:

- ☒ a. The meaning of valid-invalid bit in page table is different in paging and demand-paging. ✓
- ☒ b. Demand paging requires additional hardware support, compared to paging. ✓
- ☐ c. Paging requires some hardware support in CPU
- ☒ d. With paging, it's possible to have user programs bigger than physical memory. ✗
- ☒ e. Both demand paging and paging support shared memory pages. ✓
- ☐ f. Demand paging always increases effective memory access time.
- ☒ g. With demand paging, it's possible to have user programs bigger than physical memory. ✓
- ☒ h. Calculations of number of bits for page number and offset are same in paging and demand paging. ✓
- ☐ i. TLB hit ration has zero impact in effective memory access time in demand paging.
- ☐ j. Paging requires NO hardware support in CPU

Your answer is partially correct.

You have correctly selected 5.

The correct answers are: Demand paging requires additional hardware support, compared to paging., Both demand paging and paging support shared memory pages., With demand paging, it's possible to have user programs bigger than physical memory., Demand paging always increases effective memory access time., Paging requires some hardware support in CPU, Calculations of number of bits for page number and offset are same in paging and demand paging., The meaning of valid-invalid bit in page table is different in paging and demand-paging.

Question 4

Partially correct

Mark 0.44 out of 0.50

Suppose a processor supports base(relocation register) + limit scheme of MMU.

Assuming this, mark the statements as True/False

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	The OS may terminate the process while handling the interrupt of memory violation	✓
<input checked="" type="radio"/>	<input type="radio"/>	The hardware detects any memory access beyond the limit value and raises an interrupt	✓
<input type="radio"/>	<input checked="" type="radio"/>	The hardware may terminate the process while handling the interrupt of memory violation	✗
<input checked="" type="radio"/>	<input type="radio"/>	The OS sets up the relocation and limit registers when the process is scheduled	✓
<input checked="" type="radio"/>	<input type="radio"/>	The compiler generates machine code assuming continuous memory address space for process, and calculating appropriate sizes for code, and data;	✓
<input type="radio"/>	<input checked="" type="radio"/>	The process sets up it's own relocation and limit registers when the process is scheduled	✓
<input type="radio"/>	<input checked="" type="radio"/>	The OS detects any memory access beyond the limit value and raises an interrupt	✓
<input type="radio"/>	<input checked="" type="radio"/>	The compiler generates machine code assuming appropriately sized segments for code, data and stack.	✓

The OS may terminate the process while handling the interrupt of memory violation: True

The hardware detects any memory access beyond the limit value and raises an interrupt: True

The hardware may terminate the process while handling the interrupt of memory violation: False

The OS sets up the relocation and limit registers when the process is scheduled: True

The compiler generates machine code assuming continuous memory address space for process, and calculating appropriate sizes for code, and data;: True

The process sets up it's own relocation and limit registers when the process is scheduled: False

The OS detects any memory access beyond the limit value and raises an interrupt: False

The compiler generates machine code assuming appropriately sized segments for code, data and stack.: False

Question 5

Correct

Mark 0.50 out of 0.50

Consider the following list of free chunks, in continuous memory management:

10k, 25k, 12k, 7k, 9k, 13k

Suppose there is a request for chunk of size 9k, then the free chunk selected under each of the following schemes will be

Best fit:

9k



First fit:

10k



Worst fit:

25k



Question 6

Partially correct

Mark 0.50 out of 1.00

Select all the correct statements about MMU and it's functionality

Select one or more:

- ☐ a. MMU is a separate chip outside the processor
- ☒ b. MMU is inside the processor
- ☐ c. Logical to physical address translations in MMU are done with specific machine instructions
- ☒ d. The operating system interacts with MMU for every single address translation
- ☒ e. Illegal memory access is detected in hardware by MMU and a trap is raised
- ☐ f. The Operating system sets up relevant CPU registers to enable proper MMU translations
- ☒ g. Logical to physical address translations in MMU are done in hardware, automatically
- ☐ h. Illegal memory access is detected by operating system



Your answer is partially correct.

You have correctly selected 3.

The correct answers are: MMU is inside the processor, Logical to physical address translations in MMU are done in hardware, automatically, The Operating system sets up relevant CPU registers to enable proper MMU translations, Illegal memory access is detected in hardware by MMU and a trap is raised

Question 7

Incorrect

Mark 0.00 out of 0.50

Assuming a 8- KB page size, what is the page numbers for the address 874815 reference in decimal :
(give answer also in decimal)

Answer: ✖

The correct answer is: 107

Question 8

Incorrect

Mark 0.00 out of 0.25

Select the compiler's view of the process's address space, for each of the following MMU schemes:
(Assume that each scheme,e.g. paging/segmentation/etc is effectively utilised)

Segmentation, then paging	Many continuous chunks each of page size	✖
Relocation + Limit	Many continuous chunks of same size	✖
Segmentation	one continuous chunk	✖
Paging	many continuous chunks of variable size	✖

Your answer is incorrect.

The correct answer is: Segmentation, then paging → many continuous chunks of variable size, Relocation + Limit → one continuous chunk, Segmentation → many continuous chunks of variable size, Paging → one continuous chunk

Question 9

Incorrect

Mark 0.00 out of 0.50

Suppose the memory access time is 180ns and TLB hit ratio is 0.3, then effective memory access time is (in nanoseconds);

Answer: ✖

The correct answer is: 306.00

Question **10**

Correct

Mark 0.50 out of 0.50

In xv6, The struct context is given as

```
struct context {
    uint edi;
    uint esi;
    uint ebx;
    uint ebp;
    uint eip;
};
```

Select all the reasons that explain why only these 5 registers are included in the struct context.

- ☒ a. The segment registers are same across all contexts, hence they need not be saved ✓
- ☒ b. esp is not saved in context, because context{} is on stack and it's address is always argument to switch() ✓
- ☐ c. xv6 tries to minimize the size of context to save memory space
- ☐ d. esp is not saved in context, because it's not part of the context
- ☒ e. eax, ecx, edx are caller save, hence no need to save ✓

Your answer is correct.

The correct answers are: The segment registers are same across all contexts, hence they need not be saved, eax, ecx, edx are caller save, hence no need to save, esp is not saved in context, because context{} is on stack and it's address is always argument to switch()

Question **11**

Partially correct

Mark 0.83 out of 1.50

Arrange the following events in order, in page fault handling:

Disk interrupt wakes up the process	7	✓
The reference bit is found to be invalid by MMU	1	✓
OS makes available an empty frame	6	✗
Restart the instruction that caused the page fault	9	✓
A hardware interrupt is issued	3	✗
OS schedules a disk read for the page (from backing store)	5	✓
Process is kept in wait state	4	✗
Page tables are updated for the process	8	✓
Operating system decides that the page was not in memory	2	✗

Your answer is partially correct.

You have correctly selected 5.

The correct answer is: Disk interrupt wakes up the process → 7, The reference bit is found to be invalid by MMU → 1, OS makes available an empty frame → 4, Restart the instruction that caused the page fault → 9, A hardware interrupt is issued → 2, OS schedules a disk read for the page (from backing store) → 5, Process is kept in wait state → 6, Page tables are updated for the process → 8, Operating system decides that the page was not in memory → 3

Question **12**

Incorrect

Mark 0.00 out of 0.50

Suppose a kernel uses a buddy allocator. The smallest chunk that can be allocated is of size 32 bytes. One bit is used to track each such chunk, where 1 means allocated and 0 means free. The chunk looks like this as of now:

00001010

Now, there is a request for a chunk of 70 bytes.

After this allocation, the bitmap, indicating the status of the buddy allocator will be

Answer: 11101010



The correct answer is: 11111010

Question **13**

Incorrect

Mark 0.00 out of 0.25

The complete range of virtual addresses (after main() in main.c is over), from which the free pages used by kalloc() and kfree() is derived, are:

- ☐ a. end, 4MB
- ☐ b. P2V(end), P2V(PHYSTOP)
- ☐ c. end, P2V(4MB + PHYSTOP)
- ☒ d. P2V(end), PHYSTOP
- ☐ e. end, (4MB + PHYSTOP)
- ☐ f. end, PHYSTOP
- ☐ g. end, P2V(PHYSTOP)



Your answer is incorrect.

The correct answer is: end, P2V(PHYSTOP)

Question 14

Partially correct

Mark 0.33 out of 0.50

Match the pair

Hashed page table	Linear search on collision done by OS (e.g. SPARC Solaris) typically	✓
Inverted Page table	Linear/Parallel search using frame number in page table	✗
Hierarchical Paging	More memory access time per hierarchy	✓

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: Hashed page table → Linear search on collision done by OS (e.g. SPARC Solaris) typically, Inverted Page table → Linear/Parallel search using page number in page table, Hierarchical Paging → More memory access time per hierarchy

Question 15

Partially correct

Mark 0.29 out of 0.50

After virtual memory is implemented

(select T/F for each of the following)One Program's size can be larger than physical memory size

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	Code need not be completely in memory	✓
<input checked="" type="radio"/>	<input type="radio"/>	Cumulative size of all programs can be larger than physical memory size	✓
<input type="radio"/>	<input checked="" type="radio"/>	Virtual access to memory is granted	✗
<input checked="" type="radio"/>	<input type="radio"/>	Logical address space could be larger than physical address space	✓
<input type="radio"/>	<input checked="" type="radio"/>	Virtual addresses are available	✗
<input type="radio"/>	<input checked="" type="radio"/>	Relatively less I/O may be possible during process execution	✗
<input checked="" type="radio"/>	<input type="radio"/>	One Program's size can be larger than physical memory size	✓

Code need not be completely in memory: True

Cumulative size of all programs can be larger than physical memory size: True

Virtual access to memory is granted: False

Logical address space could be larger than physical address space: True

Virtual addresses are available: False

Relatively less I/O may be possible during process execution: True

One Program's size can be larger than physical memory size: True

Question **16**

Partially correct

Mark 0.64 out of 1.00

W.r.t. Memory management in xv6,

xv6 uses physical memory upto 224 MB only Mark statements True or False

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	The switchkvm() call in scheduler() is invoked after control comes to it from sched(), thus demanding execution in kernel's context	✓
<input checked="" type="radio"/>	<input type="radio"/>	The stack allocated in entry.S is used as stack for scheduler's context for first processor	✓
<input checked="" type="radio"/>	<input type="radio"/>	The switchkvm() call in scheduler() changes CR3 to use page directory kpgdir	✓
<input type="radio"/>	<input checked="" type="radio"/>	The free page-frame are created out of nearly 222 MB	✗
<input checked="" type="radio"/>	<input type="radio"/>	The kernel code and data take up less than 2 MB space	✓
<input type="radio"/>	<input checked="" type="radio"/>	The switchkvm() call in scheduler() changes CR3 to use page directory of new process	✗
<input type="radio"/>	<input checked="" type="radio"/>	The switchkvm() call in scheduler() is invoked after control comes to it from swtch() scheduler(), thus demanding execution in new process's context	✓
<input checked="" type="radio"/>	<input type="radio"/>	PHYSTOP can be increased to some extent, simply by editing memlayout.h	✓
<input type="radio"/>	<input checked="" type="radio"/>	xv6 uses physical memory upto 224 MB only	✗
<input checked="" type="radio"/>	<input type="radio"/>	The process's address space gets mapped on frames, obtained from ~2MB:224MB range	✓
<input type="radio"/>	<input checked="" type="radio"/>	The kernel's page table given by kpgdir variable is used as stack for scheduler's context	✗

The switchkvm() call in scheduler() is invoked after control comes to it from sched(), thus demanding execution in kernel's context: True

The stack allocated in entry.S is used as stack for scheduler's context for first processor: True

The switchkvm() call in scheduler() changes CR3 to use page directory kpgdir: True

The free page-frame are created out of nearly 222 MB: True

The kernel code and data take up less than 2 MB space: True

The switchkvm() call in scheduler() changes CR3 to use page directory of new process: False

The switchkvm() call in scheduler() is invoked after control comes to it from swtch() scheduler(), thus demanding execution in new process's context: False

PHYSTOP can be increased to some extent, simply by editing memlayout.h: True

xv6 uses physical memory upto 224 MB only: True

The process's address space gets mapped on frames, obtained from ~2MB:224MB range: True

The kernel's page table given by kpgdir variable is used as stack for scheduler's context: False

Question **17**


Incorrect

Mark 0.00 out of 1.50

Consider the reference string

6 4 2 0 1 2 6 9 2 0 5

If the number of page frames is 3, then total number of page faults (including initial), using LRU replacement is:

Answer: 

#6# 6,4# 6,4,2 # 0,4,2#0,1,2#6,1,2#6,9,2#0,9,2#0,5,2

The correct answer is: 9

Question 18

Partially correct

Mark 0.31 out of 0.50

Consider the image given below, which explains how paging works.




Figure 9.8 Paging hardware.

Mention whether each statement is True or False, with respect to this image.

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	The PTBR is present in the CPU as a register	✓
<input type="radio"/>	<input checked="" type="radio"/>	The page table is indexed using frame number	✓
<input checked="" type="radio"/>	<input type="radio"/>	The page table is indexed using page number	✗
<input checked="" type="radio"/>	<input type="radio"/>	The locating of the page table using PTBR also involves paging translation	✗
<input type="radio"/>	<input checked="" type="radio"/>	Size of page table is always determined by the size of RAM	✓
<input checked="" type="radio"/>	<input type="radio"/>	The page table is itself present in Physical memory	✓
<input checked="" type="radio"/>	<input type="radio"/>	Maximum Size of page table is determined by number of bits used for page number	✗
<input checked="" type="radio"/>	<input type="radio"/>	The physical address may not be of the same size (in bits) as the logical address	✓

The PTBR is present in the CPU as a register: True

The page table is indexed using frame number: False

The page table is indexed using page number: True

The locating of the page table using PTBR also involves paging translation: False

Size of page table is always determined by the size of RAM: False

The page table is itself present in Physical memory: True

Maximum Size of page table is determined by number of bits used for page number: True

The physical address may not be of the same size (in bits) as the logical address: True

Question **19**

Correct

Mark 2.00 out of 2.00

Given below is shared memory code with two processes sharing a memory segment.

The first process sends a user input string to second process. The second capitalizes the string. Then the first process prints the capitalized version.

Fill in the blanks to complete the code.

// First process

#define SHMSZ 27

int main()

{

char c;

int shmid;

key_t key;

char *shm, *s, string[128];

key = 5679;

if ((shmid =

shmget

✓ (key, SHMSZ, IPC_CREAT | 0666) < 0) {

perror("shmget");

exit(1);

}

if ((shm =

shmat

✓ (shmid, NULL, 0)) == (char *) -1) {

perror("shmat");

exit(1);

}

s = shm;

*s = '\$';

scanf("%s", string);

strcpy(s + 1, string);

*s = '

@

✓ '; //note the quotes

while(*s != ')

\$

✓ ')

sleep(1);

printf("%s\n", s + 1);

exit(0);

}

//Second process

#define SHMSZ 27

int main()

{

int shmid;

key_t key;

char *shm, *s;

int i;

char string[128];

key =

5679

```

✓ ;
if ((shm = shmget(key, SHMSZ, 0666)) < 0) {
    perror("shmget");
    exit(1);
}
if ((shm = shmat(shmid, NULL, 0)) == (char *) -1) {
    perror("shmat");
    exit(1);
}
s =

```

shm

```

✓ ;
while(*s != '@')
    sleep(1);
for(i = 0; i < strlen(s + 1); i++)
    s[i + 1] = toupper(s[i + 1]);
*s = '$';
exit(0);
}

```

Question 20

Partially correct

Mark 0.25 out of 0.50

Map the functionality/use with function/variable in xv6 code.

return a free page, if available; 0, otherwise

kinit1()

✗

Create page table entries for a given range of virtual and physical addresses; including page directory entries if needed

mappages()

✓

Array listing the kernel memory mappings, to be used by setupkvm()

kmap[]

✓

Setup kernel part of a page table, mapping kernel code, data, read-only data, I/O space, devices

kvmalloc()

✗

Return address of page table entry in a given page directory, for a given virtual address; creates page table if necessary

walkpgdir()

✓

Setup kernel part of a page table, and switch to that page table

setupkvm()

✗

Your answer is partially correct.

You have correctly selected 3.

The correct answer is: return a free page, if available; 0, otherwise → kalloc(), Create page table entries for a given range of virtual and physical addresses; including page directory entries if needed → mappages(), Array listing the kernel memory mappings, to be used by setupkvm() → kmap[], Setup kernel part of a page table, mapping kernel code, data, read-only data, I/O space, devices → setupkvm(), Return address of page table entry in a given page directory, for a given virtual address; creates page table if necessary → walkpgdir(), Setup kernel part of a page table, and switch to that page table → kvmalloc()

Question **21**

Partially correct

Mark 1.53 out of 2.50

Order events in xv6 timer interrupt code

(Transition from process P1 to P2's code.)

P2 is selected and marked RUNNING	12	✓
Change of stack from user stack to kernel stack of P1	3	✓
Timer interrupt occurs	2	✓
alltraps() will call iret	17	✗
change to context of P2, P2's kernel stack in use now	13	✓
P2's trap() will return to alltraps	16	✗
jump in vector.S	4	✓
P2 will return from sched() in yield()	14	✗
yield() is called	8	✓
trap() is called	7	✓
Process P2 is executing	18	✗
P1 is marked as RUNNABLE	9	✓
P2's yield() will return in trap()	15	✗
Process P1 is executing	1	✓
sched() is called,	11	✗
change to context of the scheduler, scheduler's stack in use now	10	✗
jump to alltraps	5	✓
Trapframe is built on kernel stack of P1	6	✓

Your answer is partially correct.

You have correctly selected 11.

The correct answer is: P2 is selected and marked RUNNING → 12, Change of stack from user stack to kernel stack of P1 → 3, Timer interrupt occurs → 2, alltraps() will call iret → 18, change to context of P2, P2's kernel stack in use now → 13, P2's trap() will return to alltraps → 17, jump in vector.S → 4, P2 will return from sched() in yield() → 15, yield() is called → 8, trap() is called → 7, Process P2 is executing → 14, P1 is marked as RUNNABLE → 9, P2's yield() will return in trap() → 16, Process P1 is executing → 1, sched() is called, → 10, change to context of the scheduler, scheduler's stack in use now → 11, jump to alltraps → 5, Trapframe is built on kernel stack of P1 → 6

Question **22**

Incorrect

Mark 0.00 out of 1.00

Given that the memory access time is 200 ns, probability of a page fault is 0.7 and page fault handling time is 8 ms,
The effective memory access time in nanoseconds is:

Answer: ❌

The correct answer is: 5600060.00

Question **23**

Correct

Mark 0.25 out of 0.25

Select the state that is not possible after the given state, for a process:

New: ✔️

Ready : ✔️

Running: : ✔️

Waiting: ✔️

Question **24**

Partially correct

Mark 0.63 out of 1.00

Select the correct statements about sched() and scheduler() in xv6 code

- ☒ a. scheduler() switches to the selected process's context ✔️
- ☒ b. When either sched() or scheduler() is called, it does not return immediately to caller ✔️
- ☐ c. After call to switch() in sched(), the control moves to code in scheduler()
- ☒ d. Each call to sched() or scheduler() involves change of one stack inside switch() ✔️
- ☐ e. After call to switch() in scheduler(), the control moves to code in sched()
- ☒ f. When either sched() or scheduler() is called, it results in a context switch ✔️
- ☒ g. sched() switches to the scheduler's context ✔️
- ☐ h. sched() and scheduler() are co-routines

Your answer is partially correct.

You have correctly selected 5.

The correct answers are: sched() and scheduler() are co-routines, When either sched() or scheduler() is called, it does not return immediately to caller, When either sched() or scheduler() is called, it results in a context switch, sched() switches to the scheduler's context, scheduler() switches to the selected process's context, After call to switch() in scheduler(), the control moves to code in sched(), After call to switch() in sched(), the control moves to code in scheduler(), Each call to sched() or scheduler() involves change of one stack inside switch()

Question **25**

Correct

Mark 0.25 out of 0.25

The data structure used in `kalloc()` and `kfree()` in `xv6` is

- ☐ a. Doubly linked circular list
- ☐ b. Singly linked circular list
- ☐ c. Double linked NULL terminated list
- ☒ d. Singly linked NULL terminated list



Your answer is correct.

The correct answer is: Singly linked NULL terminated list

[◀ \(Assignment\) lseek system call in xv6](#)

Jump to...

[Dashboard](#) / [My courses](#) / [Computer Engineering & IT](#) / [CEIT-Even-sem-20-21](#) / [OS-Even-sem-2020-21](#) / [16 May - 22 May](#) / [End Sem Exam OS-2021](#)

Started on Saturday, 22 May 2021, 8:00 AM

State Finished

Completed on Saturday, 22 May 2021, 9:30 AM

Time taken 1 hour 30 mins

Grade 26.12 out of 40.00 (65%)

Question 1

Incorrect

Mark 0.00 out of 1.00

A 4 GB disk with 1 KB of block size would require these many number of **blocks** for it's free block bitmap:

Answer: ✖

The correct answer is: 512

Question 2

Correct

Mark 1.00 out of 1.00

Given that the memory access time is 110 ns, probability of a page fault is 0.5 and page fault handling time is 12 ms,
The effective memory access time in nanoseconds is:

Answer: ✔

The correct answer is: 6000055.00

Question 3

Incorrect

Mark 0.00 out of 1.00

The maximum size of a file in number of blocks of BSIZE in xv6 code is
(write a number only)

Answer: ✖

The correct answer is: 138

Question 4

Incorrect

Mark 0.00 out of 1.00

Calculate the average waiting time using
Round Robin scheduling with time quantum of 5 time units
for the following workload

assuming that they arrive in the order written below.

Process Burst Time

P1 5

P2 7

P3 6

P4 2

Write only a number in the answer upto two decimal points.

Answer: ✖

The correct answer is: 10.25



Question 5

Correct

Mark 1.00 out of 1.00

For the reference string

4 2 5 1 0 1 2 5 4 1 2

the number of page faults, including initial ones,
with FIFO replacement and 2 frames are :

Answer: ✓

4 -

4 2

5 2

5 1

0 1

-

2 1

2 5

4 5

4 1

2 1

The correct answer is: 10

Question 6

Correct

Mark 1.00 out of 1.00

Assuming a 16- KB page size, what is the page number for the address 428517 reference in decimal :
(give answer also in decimal)

Answer: ✓

The correct answer is: 26

Question 7

Correct

Mark 1.00 out of 1.00

In the code below assume that each function can be executed concurrently by many threads/processes.
Ignore syntactical issues, and focus on the semantics.

This program is an example of

```
spinlock a, b; // assume initialized
thread1() {
    spinlock(b);
    //some code;
    spinlock(a);
    //some code;
    spinunlock(b);
    spinunlock(a);
}
thread2() {
    spinlock(a);
    //some code;
    spinlock(b);
    //some code;
    spinunlock(b);
    spinunlock(a);
}
```

- ☒ a. Deadlock
- ☐ b. Self Deadlock
- ☐ c. None of these
- ☐ d. Deadlock or livelock depending on actual race
- ☐ e. Livelock



Your answer is correct.

The correct answer is: Deadlock



Question 8

Partially correct

Mark 1.33 out of 2.00

Match the snippets of xv6 code with the core functionality they achieve, or problems they avoid.
 "... means some code.

```
static inline uint
xchg(volatile uint *addr, uint newval)
{
    uint result;
```

// The + in "+m" denotes a read-modify-write operand.

```
asm volatile("lock; xchgl %0, %1" :
    "+m" (*addr), "=a" (result) :
    "1" (newval) :
    "cc");
return result;
}
```

Atomic compare and swap instruction (to be expanded inline into code)



```
void
sleep(void *chan, struct spinlock *lk)
{
    ...
    if(lk != &ptable.lock){
        acquire(&ptable.lock);
        release(lk);
    }
```

If you don't do this, a process may be running on two processors parallely



```
void
acquire(struct spinlock *lk)
{
    ...
    __sync_synchronize();
```

Tell compiler not to reorder memory access beyond this line



Your answer is partially correct.

You have correctly selected 2.

The correct answer is: static inline uint

```
xchg(volatile uint *addr, uint newval)
{
    uint result;
```

// The + in "+m" denotes a read-modify-write operand.

```
asm volatile("lock; xchgl %0, %1" :
    "+m" (*addr), "=a" (result) :
    "1" (newval) :
    "cc");
return result;
```

} → Atomic compare and swap instruction (to be expanded inline into code), void

```
sleep(void *chan, struct spinlock *lk)
```

```
{
    ...
    if(lk != &ptable.lock){
        acquire(&ptable.lock);
        release(lk);
```

} → Avoid a self-deadlock, void

```
acquire(struct spinlock *lk)
```

```
{
```

```
...
```

__sync_synchronize(); → Tell compiler not to reorder memory access beyond this line

Question 9

Correct

Mark 1.00 out of 1.00

Predict the output of the program given here.

Assume that all the path names for the programs are correct. For example "/usr/bin/echo" will actually run echo command.

Assume that there is no mixing of printf output on screen if two of them run concurrently.

In the answer replace a new line by a single space.

For example::

good

output

should be written as good output

--

```
main() {
    int i;
    i = fork();
    if(i == 0)
        execl("/usr/bin/echo", "/usr/bin/echo", "hi", 0);
    else
        wait(0);
    fork();
    execl("/usr/bin/echo", "/usr/bin/echo", "one", 0);
}
```

Answer: hi one one



The correct answer is: hi one one

Question 10

Partially correct

Mark 1.67 out of 2.00

Select all the blocks that may need to be written back to disk (if updated, of-course), as "Yes", when an operation of deleting a file is carried out on ext2 file system.

An option has to be correct entirely to be marked "Yes"

Superblock

Yes



One or multiple data blocks of the parent directory

No



One or more data bitmap blocks for the parent directory

No



Block bitmap(s) for all the blocks of the file

No



Possibly one block bitmap corresponding to the parent directory

Yes



Data blocks of the file

No



Your answer is partially correct.

only one data block of parent directory. multiple blocks not possible. an entry is always contained within one single block

You have correctly selected 5.

The correct answer is: Superblock → Yes, One or multiple data blocks of the parent directory → No, One or more data bitmap blocks for the parent directory → No, Block bitmap(s) for all the blocks of the file → Yes, Possibly one block bitmap corresponding to the parent directory → Yes, Data blocks of the file → No



Question 11

Correct

Mark 1.00 out of 1.00

Select all the correct statements about bootloader.

Every wrong selection will deduct marks proportional to $1/n$ where n is total wrong choices in the question.

You will get minimum a zero.

- ☒ a. Modern Bootloaders often allow configuring the way an OS boots
- ☒ b. Bootloaders allow selection of OS to boot from
- ☐ c. Bootloader must be one sector in length
- ☐ d. The bootloader loads the BIOS
- ☒ e. LILO is a bootloader



Your answer is correct.

The correct answers are: LILO is a bootloader, Modern Bootloaders often allow configuring the way an OS boots, Bootloaders allow selection of OS to boot from



Question 12




Incorrect

Mark 0.00 out of 1.00


For each of the resource allocation diagram shown,
infer whether the graph contains at least one deadlock or not.

Yes

No



: Yes



Question 13

Partially correct

Mark 0.71 out of 1.00

Mark the statements about device drivers by marking as True or False.

True	False		
<input type="radio"/>	<input checked="" type="radio"/>	It's possible that a particular hardware has multiple device drivers available for it.	✗
<input checked="" type="radio"/>	<input type="radio"/>	xv6 has device drivers for IDE disk and console.	✓
<input checked="" type="radio"/>	<input type="radio"/>	A disk driver converts OS's logical view of disk into physical locations on disk.	✓
<input checked="" type="radio"/>	<input type="radio"/>	A device driver code is specific to a hardware device	✓
<input checked="" type="radio"/>	<input type="radio"/>	All devices of the same type (e.g. 2 hard disks) can typically use the same device driver	✓
<input type="radio"/>	<input checked="" type="radio"/>	Writing a device driver mandatorily demands reading the technical documentation about the hardware.	✗
<input type="radio"/>	<input checked="" type="radio"/>	Device driver is an intermediary between the end-user and OS	✓

It's possible that a particular hardware has multiple device drivers available for it.: True

xv6 has device drivers for IDE disk and console.: True

A disk driver converts OS's logical view of disk into physical locations on disk.: True

A device driver code is specific to a hardware device: True

All devices of the same type (e.g. 2 hard disks) can typically use the same device driver: True

Writing a device driver mandatorily demands reading the technical documentation about the hardware.: True

Device driver is an intermediary between the end-user and OS: False

Question 14

Partially correct

Mark 0.33 out of 1.00

Consider this program.

Some statements are identified using the // comment at the end.

Assume that = is an atomic operation.

```
#include <stdio.h>
#include <pthread.h>
long c = 0, c1 = 0, c2 = 0, run = 1;
void *thread1(void *arg) {
    while(run == 1) { //E
        c = 10; //A
        c1 = c2 + 5; //B
    }
}
void *thread2(void *arg) {
    while(run == 1) { //F
        c = 20; //C
        c2 = c1 + 3; //D
    }
}
int main() {
    pthread_t th1, th2;
    pthread_create(&th1, NULL, thread1, NULL);
    pthread_create(&th2, NULL, thread2, NULL);
    sleep(2);
    run = 0;
    fprintf(stdout, "c = %ld c1+c2 = %ld c1 = %ld c2 = %ld \n", c, c1+c2, c1, c2);
    fflush(stdout);
}
```

Which statements are part of the critical Section?

Yes	No	
<input checked="" type="radio"/>	<input checked="" type="radio"/>	F
<input checked="" type="radio"/>	<input type="radio"/>	D
<input checked="" type="radio"/>	<input checked="" type="radio"/>	C
<input checked="" type="radio"/>	<input checked="" type="radio"/>	A
<input checked="" type="radio"/>	<input type="radio"/>	B
<input checked="" type="radio"/>	<input checked="" type="radio"/>	E

F: No

D: Yes

C: No

A: No

B: Yes

E: No

Question 15

Partially correct

Mark 1.43 out of 2.00

Mark statements as T/F

All statements are in the context of preventing deadlocks.

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	A process holding one resources and waiting for just one more resource can also be involved in a deadlock.	✓
<input type="radio"/>	<input checked="" type="radio"/>	If a resource allocation graph contains a cycle then there is a guarantee of a deadlock	✗
<input type="radio"/>	<input checked="" type="radio"/>	The lock ordering to be followed to avoid circular wait is a code in OS that checks for compliance with decided order	✗
<input checked="" type="radio"/>	<input type="radio"/>	Circular wait is avoided by enforcing a lock ordering	✓
<input checked="" type="radio"/>	<input type="radio"/>	Hold and wait means a thread/process holding some locks and waiting for acquiring some.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Deadlock is possible if all the conditions are met at the same time: Mutual exclusion, hold and wait, no pre-emption, circular wait.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Mutual exclusion is a necessary condition for deadlock because it brings in locks on which deadlock happens	✓

A process holding one resources and waiting for just one more resource can also be involved in a deadlock.: True

If a resource allocation graph contains a cycle then there is a guarantee of a deadlock: False

The lock ordering to be followed to avoid circular wait is a code in OS that checks for compliance with decided order: False

Circular wait is avoided by enforcing a lock ordering: True

Hold and wait means a thread/process holding some locks and waiting for acquiring some.: True

Deadlock is possible if all the conditions are met at the same time: Mutual exclusion, hold and wait, no pre-emption, circular wait.: True

Mutual exclusion is a necessary condition for deadlock because it brings in locks on which deadlock happens: True

Question 16

Correct

Mark 1.00 out of 1.00

Match the left side use(or non-use) of a synchronization primitive with the best option on the right side.

This is the smallest primitive made available in software, using the hardware provided atomic instructions	spinlock	✓
This tool is useful for event-wait scenarios	semaphore	✓
This tool is more useful on multiprocessor systems	spinlock	✓
This tool is quite attractive in solving the main bounded buffer problem	semaphore	✓
This tool is very useful for waiting for 'something'	condition variables	✓

Your answer is correct.

The correct answer is: This is the smallest primitive made available in software, using the hardware provided atomic instructions → spinlock, This tool is useful for event-wait scenarios → semaphore, This tool is more useful on multiprocessor systems → spinlock, This tool is quite attractive in solving the main bounded buffer problem → semaphore, This tool is very useful for waiting for 'something' → condition variables

Question 17

Correct

Mark 1.00 out of 1.00

The permissions -rwx--x--x on a file mean

- ☒ a. The file can be read only by the owner
- ☐ b. 'cat' on the file by owner will not work
- ☐ c. 'cat' on the file by any user will work
- ☒ d. 'rm' on the file by any user will work
- ☒ e. The file can be executed by anyone
- ☒ f. The file can be written only by the owner

✓

✓

✓

✓

Your answer is correct.

The correct answers are: The file can be executed by anyone, The file can be read only by the owner, The file can be written only by the owner, 'rm' on the file by any user will work

Question 18

Incorrect

Mark 0.00 out of 1.00

Note: for this question you get full marks if you select all and only correct options, you get ZERO if at least one option is wrong or not selected.

Select all the correct statements about log structured file systems.

- ☒ a. a transaction is said to be committed when all operations are written to file system
- ☒ b. log may be kept on same block device or another block device
- ☐ c. file system recovery may end up losing data
- ☒ d. even if file systems followed immediate writes (i.e. non-delayed writes), it could still require recovery
- ☒ e. file system recovery recovers all the lost data

✗

✓

✓

✗

Your answer is incorrect.

The correct answers are: file system recovery may end up losing data, log may be kept on same block device or another block device, even if file systems followed immediate writes (i.e. non-delayed writes), it could still require recovery

Question 19

Incorrect

Mark 0.00 out of 1.00

Consider the structure of directory entry in ext2, as shown in this diagram.

	inode	rec_len	file_type	name_len	name
0	21	12	1	2	· \0 \0 \0
12	22	12	2	2	· · \0 \0
24	53	16	5	2	h o m e 1 \0 \0 \0
40	67	28	3	2	u s r \0
52	0	16	7	1	o l d f i 1 e \0
68	34	12	4	2	s b i n

Select the correct statements about the directory entry in ext2 file system.

The correct formula for rec_len is (when entries are continuously stored)

- ☐ a. $\text{rec_len} = \text{sizeof}(\text{inode entry}) + \text{sizeof}(\text{name len entry}) + \text{sizeof}(\text{file type entry}) + (\text{strlen}(\text{name}) + (-1) * (\text{strlen}(\text{name}) \% 4)$
- ☐ b. $\text{rec_len} = \text{sizeof}(\text{inode entry}) + \text{sizeof}(\text{name len entry}) + \text{sizeof}(\text{file type entry}) + (\text{strlen}(\text{name}) + (\text{strlen}(\text{name}) - 4) \% 4$
- ☒ c. $\text{rec_len} = \text{sizeof}(\text{inode entry}) + \text{sizeof}(\text{name len entry}) + \text{sizeof}(\text{file type entry}) + (\text{strlen}(\text{name}) + 4 - (\text{strlen}(\text{name}) \% 4)$
- ☐ d. $\text{rec_len} = \text{sizeof}(\text{inode entry}) + \text{sizeof}(\text{name len entry}) + \text{sizeof}(\text{file type entry}) + (\text{strlen}(\text{name}) + (-1) * (\text{strlen}(\text{name}) - 4)$
- ☐ e. $\text{rec_len} = \text{sizeof}(\text{inode entry}) + \text{sizeof}(\text{name len entry}) + \text{sizeof}(\text{file type entry}) + (\text{strlen}(\text{name}) \% 4$
- ☐ f. $\text{rec_len} = \text{sizeof}(\text{inode entry}) + \text{sizeof}(\text{name len entry}) + \text{sizeof}(\text{file type entry}) + \text{strlen}(\text{name})$

Your answer is incorrect.

The correct answer is: $\text{rec_len} = \text{sizeof}(\text{inode entry}) + \text{sizeof}(\text{name len entry}) + \text{sizeof}(\text{file type entry}) + (\text{strlen}(\text{name}) + (-1) * (\text{strlen}(\text{name}) - 4)$

Question 20

Partially correct

Mark 0.50 out of 1.00

Mark whether the given sequence of events is possible or not-possible. Also, select the reason for your answer.

For each sequence it's a not-possible sequence if some important event is not mentioned in the sequence.

Assume that the kernel code is non-interruptible and uniprocessor system.

Process P1 executing a system call

Timer interrupt

Generic interrupt handler runs

Scheduler runs

Scheduler selects P2 for execution

P2 returns from timer interrupt handler

Process p2, user code executing

This sequence of events is: ✓

Because

✗

Question 21

Incorrect

Mark 0.00 out of 1.00

The given semaphore implementation faces which problem?

Assume any suitable code for signal()

Note: blocks means waits in a wait queue.

```
struct semaphore {
    int val;
    spinlock lk;
};
sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->sl = 0;
}
wait(semaphore *s) {
    spinlock(&(s->sl));
    while(s->val <= 0)
        ;
    (s->val)--;
    spinunlock(&(s->sl));
}
```

- ☐ a. blocks holding a spinlock
- ☐ b. deadlock
- ☒ c. too much spinning, bounded wait not guaranteed
- ☐ d. not holding lock after unblock

✖

Your answer is incorrect.

The correct answer is: deadlock

Question 22

Partially correct

Mark 0.80 out of 1.00

Mark statements True/False w.r.t. change of states of a process.

Reference: The process state diagram (and your understanding of how kernel code works)




Figure 3.2 Diagram of process state.

True	False		
<input type="radio"/>	<input checked="" type="radio"/>	A process in RUNNING state only can become TERMINATED because scheduler moves it to ZOMBIE state	✓
<input checked="" type="radio"/>	<input type="radio"/>	A process in READY state can not go to WAITING state because the resource on which it will WAIT will not be in use when process is in READY state.	✗
<input checked="" type="radio"/>	<input type="radio"/>	A process in WAITING state can not become RUNNING because the event it's waiting for has not occurred	✓
<input checked="" type="radio"/>	<input type="radio"/>	Every process has to go through ZOMBIE state, at least for a small duration.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Only a process in READY state is considered by scheduler	✓

A process in RUNNING state only can become TERMINATED because scheduler moves it to ZOMBIE state: False

A process in READY state can not go to WAITING state because the resource on which it will WAIT will not be in use when process is in READY state.: False

A process in WAITING state can not become RUNNING because the event it's waiting for has not occurred: True

Every process has to go through ZOMBIE state, at least for a small duration.: True

Only a process in READY state is considered by scheduler: True

Question 23

Correct

Mark 1.00 out of 1.00

Select T/F for statements about Volume Managers.

Do pay attention to the use of the words physical partition and physical volume.

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	The volume manager can create further internal sub-divisions of a physical partition for efficiency or features.	✓
<input checked="" type="radio"/>	<input type="radio"/>	A logical volume can be extended in size but upto the size of volume group	✓
<input checked="" type="radio"/>	<input type="radio"/>	A logical volume may span across multiple physical volumes	✓
<input checked="" type="radio"/>	<input type="radio"/>	The volume manager stores additional metadata on the physical disk partitions	✓
<input checked="" type="radio"/>	<input type="radio"/>	A physical partition should be initialized as a physical volume, before it can be used by volume manager.	✓
<input checked="" type="radio"/>	<input type="radio"/>	A volume group consists of multiple physical volumes	✓
<input checked="" type="radio"/>	<input type="radio"/>	A logical volume may span across multiple physical partitions	✓ since a physical volume is made up of physical partitions, and a volume can span across multiple PVs, it can also span across multiple PP

The volume manager can create further internal sub-divisions of a physical partition for efficiency or features.: True

A logical volume can be extended in size but upto the size of volume group: True

A logical volume may span across multiple physical volumes: True

The volume manager stores additional metadata on the physical disk partitions: True

A physical partition should be initialized as a physical volume, before it can be used by volume manager.: True

A volume group consists of multiple physical volumes: True

A logical volume may span across multiple physical partitions: True

Question 24

Correct

Mark 1.00 out of 1.00

Map the block allocation scheme with the problem it suffers from

(Match pairs 1-1, match a scheme with the problem that it suffers from relatively the most, compared to others)

Continuous allocation	need for compaction	✓
Linked allocation	Too many seeks	✓
Indexed Allocation	Overhead of reading metadata blocks	✓

Your answer is correct.

The correct answer is: Continuous allocation → need for compaction, Linked allocation → Too many seeks, Indexed Allocation → Overhead of reading metadata blocks

Question 25

Correct

Mark 1.00 out of 1.00

This one is not a system call:

- ☐ a. open
☐ b. read
☐ c. write
☒ d. scheduler

✓

Your answer is correct.

The correct answer is: scheduler



Question 26

Correct

Mark 1.00 out of 1.00

Match the pairs.

This question is based on your general knowledge about operating systems/related concepts and their features.

Java threads	monitors, re-entrant locks, semaphores	✓
Linux threads	atomic-instructions, spinlocks, etc.	✓
POSIX threads	semaphore, mutex, condition variables	✓

Your answer is correct.

The correct answer is: Java threads → monitors, re-entrant locks, semaphores, Linux threads → atomic-instructions, spinlocks, etc., POSIX threads → semaphore, mutex, condition variables

Question 27

Correct

Mark 1.00 out of 1.00

Consider the following list of free chunks, in continuous memory management:

7k, 15k, 21k, 14k, 19k, 6k

Suppose there is a request for chunk of size 5k, then the free chunk selected under each of the following schemes will be

Best fit:	6k	✓
First fit:	7k	✓
Worst fit:	21k	✓

Question 28

Correct

Mark 1.00 out of 1.00

This one is not a scheduling algorithm

- ☐ a. Round Robin
- ☐ b. SJF
- ☒ c. Mergesort
- ☐ d. FCFS

✓

Your answer is correct.

The correct answer is: Mergesort

Question 29

Correct

Mark 1.00 out of 1.00

Mark whether the concept is related to scheduling or not.

Yes	No		
<input checked="" type="radio"/>	<input type="radio"/>	timer interrupt	✓
<input checked="" type="radio"/>	<input type="radio"/>	context-switch	✓
<input checked="" type="radio"/>	<input type="radio"/>	ready-queue	✓
<input type="radio"/>	<input checked="" type="radio"/>	file-table	✓
<input checked="" type="radio"/>	<input type="radio"/>	runnable process	✓

timer interrupt: Yes

context-switch: Yes

ready-queue: Yes

file-table: No

runnable process: Yes

Question 30

Partially correct

Mark 1.00 out of 2.00

Map ext2 data structure features with their purpose

Many
copies of
Superblock

Choose...

Free
blocks
count in
superblock
and group
descriptor

Redundancy to ensure the most crucial data structure is not lost

Used
directories
count in
group
descriptor

is redundant and helps do calculations of directory entries faster

Combining
file type
and
access
rights in
one
variable

saves 1 byte of space

rec_len
field in
directory
entry

Try to keep all the data of a directory and it's file close together in a group

File Name
is padded

aligns all memory accesses on word boundary, improving performance

Inode
bitmap is
one block

limits total number of files that can belong to a group

Block
bitmap is
one block

Limits the size of a block group, thus improvising on purpose of a group

Mount
count in
superblock

to enforce file check after certain amount of mounts at boot time

Inode table
location in
Group
Descriptor

is redundant and helps do calculations of directory entries faster

Inode table

All inodes are kept together so that one disk read leads to reading many inodes together, effectively doing a buffering of subsequent inode reads, and to save space on disk

A group

Redundancy to ensure the most crucial data structure is not lost

Your answer is partially correct.

You have correctly selected 6.

The correct answer is: **Many copies of Superblock** → Redundancy to ensure the most crucial data structure is not lost, **Free blocks count in superblock and group descriptor** → Redundancy to help fsck restore consistency, **Used directories count in group descriptor** → attempt is made to evenly spread the first-level directories, this count is used there, **Combining file type and access rights in one variable** → saves 1 byte of space, **rec_len field in directory entry** → allows holes and linking of entries in directory, File Name is padded → aligns all memory accesses on word boundary, improving performance, **Inode bitmap is one block** → limits total number of files that can belong to a group, **Block bitmap is one block** → Limits the size of a block group, thus improvising on purpose of a group, **Mount count in superblock** → to enforce file check after certain amount of mounts at boot time, **Inode table location in Group Descriptor** → Obvious, as it's per group and not per file-system, **Inode table** → All inodes are kept together so that one disk read leads to reading many inodes together, effectively doing a buffering of subsequent inode reads, and to save space on disk, **A group** → Try to keep all the data of a directory and it's file close together in a group



Question 31

Partially correct

Mark 1.85 out of 2.00

Mark True/False

Statements about scheduling and scheduling algorithms

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	The nice() system call is used to set priorities for processes	✓
<input checked="" type="radio"/>	<input type="radio"/>	Aging is used to ensure that low-priority processes do not starve in priority scheduling.	✓
<input type="radio"/>	<input checked="" type="radio"/>	In non-pre-emptive priority scheduling, the highest priority process is scheduled and runs until it gives up CPU.	✗
<input checked="" type="radio"/>	<input type="radio"/>	xv6 code does not care about Processor Affinity	✓
<input checked="" type="radio"/>	<input type="radio"/>	In pre-emptive priority scheduling, priority is implemented by assigning more time quantum to higher priority process.	✓
<input checked="" type="radio"/>	<input type="radio"/>	A scheduling algorithm is non-preemptive if it does context switch only if a process voluntarily relinquishes CPU or it terminates.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Processor Affinity refers to memory accesses of a process being stored on cache of that processor	✓
<input checked="" type="radio"/>	<input type="radio"/>	Response time will be quite poor on non-interruptible kernels	✓
<input checked="" type="radio"/>	<input type="radio"/>	Shortest Remaining Time First algorithm is nothing but pre-emptive Shortest Job First algorithm	✓
<input checked="" type="radio"/>	<input type="radio"/>	On Linuxes the CPU utilisation is measured as the time spent in scheduling the idle thread	✓
<input checked="" type="radio"/>	<input type="radio"/>	Generally the voluntary context switches are much more than non-voluntary context switches on a Linux system.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Pre-emptive scheduling leads to many race conditions in kernel code.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Statistical observations tell us that most processes have large number of small CPU bursts and relatively smaller numbers of large CPU bursts.	✓

The nice() system call is used to set priorities for processes: True

Aging is used to ensure that low-priority processes do not starve in priority scheduling.: True

In non-pre-emptive priority scheduling, the highest priority process is scheduled and runs until it gives up CPU.: True

xv6 code does not care about Processor Affinity: True

In pre-emptive priority scheduling, priority is implemented by assigning more time quantum to higher priority process.: True

A scheduling algorithm is non-preemptive if it does context switch only if a process voluntarily relinquishes CPU or it terminates.: True

Processor Affinity refers to memory accesses of a process being stored on cache of that processor: True

Response time will be quite poor on non-interruptible kernels: True

Shortest Remaining Time First algorithm is nothing but pre-emptive Shortest Job First algorithm: True

On Linuxes the CPU utilisation is measured as the time spent in scheduling the idle thread: True

Generally the voluntary context switches are much more than non-voluntary context switches on a Linux system.: True

Pre-emptive scheduling leads to many race conditions in kernel code.: True

Statistical observations tell us that most processes have large number of small CPU bursts and relatively smaller numbers of large CPU bursts.: True

Question 32

Partially correct

Mark 1.17 out of 2.00

The unix file semantics demand that changes to any open file are visible immediately to any other processes accessing that file at that point in time.

Select the data-structure/programmatic features that ensure the implementation of unix semantics. (Assume there is no mmap())

Yes	No		
<input type="radio"/>	<input checked="" type="radio"/>	All processes accessing the same file share the file descriptor among themselves	✓
<input type="radio"/>	<input checked="" type="radio"/>	The pointer entry in the file descriptor array entry points to the data of the file directly	✓
<input checked="" type="radio"/>	<input type="radio"/>	There is only one global file structure per on-disk file.	✗
<input type="radio"/>	<input checked="" type="radio"/>	All file accesses are made using only global variables	✓
<input checked="" type="radio"/>	<input type="radio"/>	The 'file offset' is shared among all the processes that access the file.	✗
<input type="radio"/>	<input checked="" type="radio"/>	No synchronization is implemented so that changes are made available immediately.	✓
<input checked="" type="radio"/>	<input type="radio"/>	A single spinlock is to be used to protect the unique global 'file structure' representing the file, thus synchronizing access, and making other processes wait for earlier process to finish writing so that writes get visible immediately.	✗
<input checked="" type="radio"/>	<input type="radio"/>	There is only one in-memory copy of the on disk file's contents in kernel memory/buffers	✓
<input checked="" type="radio"/>	<input type="radio"/>	The file descriptors in every PCB are pointers to the same global file structure.	✗
<input type="radio"/>	<input checked="" type="radio"/>	The file descriptor array is external to PCB and all processes that share a file, have pointers to same file-descriptors' array	✓
<input checked="" type="radio"/>	<input type="radio"/>	All file structures representing any open file, give access to the same in-memory copy of the file's contents	✓
<input checked="" type="radio"/>	<input type="radio"/>	The 'file offset' index is stored outside the file-structure to which file-descriptor array points	✗

All processes accessing the same file share the file descriptor among themselves: No

The pointer entry in the file descriptor array entry points to the data of the file directly: No

There is only one global file structure per on-disk file.: No

All file accesses are made using only global variables: No

The 'file offset' is shared among all the processes that access the file.: No

No synchronization is implemented so that changes are made available immediately.: No

A single spinlock is to be used to protect the unique global 'file structure' representing the file, thus synchronizing access, and making other processes wait for earlier process to finish writing so that writes get visible immediately.: No

There is only one in-memory copy of the on disk file's contents in kernel memory/buffers: Yes

The file descriptors in every PCB are pointers to the same global file structure.: No

The file descriptor array is external to PCB and all processes that share a file, have pointers to same file-descriptors' array: No

All file structures representing any open file, give access to the same in-memory copy of the file's contents: Yes

The 'file offset' index is stored outside the file-structure to which file-descriptor array points: No

Question 33

Partially correct

Mark 0.33 out of 2.00

Map the function in xv6's file system code, to it's perceived logical layer.

namei	inode	✗
filestat()	Choose...	
dirlookup	directory	✓
ialloc	file descriptor	✗
stati	Choose...	
ideintr	buffer cache	✗
bread	Choose...	
balloc	file descriptor	✗
sys_chdir()	system call	✓
skipelem	system call	✗
commit	system call	✗
bmap	system call	✗

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: namei → pathname lookup, filestat() → file descriptor, dirlookup → directory, ialloc → inode, stati → inode, ideintr → disk driver, bread → buffer cache, balloc → block allocation on disk, sys_chdir() → system call, skipelem → pathname lookup, commit → logging, bmap → inode

[◀ Course Exit Feedback](#)

Jump to...

[xv6-public-master ▶](#)

[Dashboard](#) / [My courses](#) / [Computer Engineering & IT](#) / [CEIT-Even-sem-21-22](#) / [OS-even-sem-21-22](#) / [24 January - 30 January](#)
/ [Topic-wise Quiz: 1 \(system calls, x86, calling convention\)](#)

Started on Monday, 24 January 2022, 7:07:42 PM

State Finished

Completed on Monday, 24 January 2022, 8:08:11 PM

Time taken 1 hour

Grade 8.90 out of 20.00 (45%)

Question **1**

Complete

Mark 0.80 out of 1.00

Match the register with the segment used with it.

ebp	<input type="text" value="ss"/>
eip	<input type="text" value="cs"/>
edi	<input type="text" value="ds"/>
esp	<input type="text" value="ss"/>
esi	<input type="text" value="ds"/>

The correct answer is: ebp → ss, eip → cs, edi → es, esp → ss, esi → ds

Question **2**

Complete

Mark 1.00 out of 1.00

```
int value = 5;
int main()
{
    pid_t pid;
    pid = fork();
    if (pid == 0) { /* child process */
        value += 15;
        return 0;
    }
    else if (pid > 0) { /* parent process */
        wait(NULL);
        printf("%d", value); /* LINE A */
    }
    return 0;
}
```

What's the value printed here at LINE A?

Answer:

The correct answer is: 5

Question 3

Complete

Mark 0.50 out of 0.50

Is the command "cat README > done &" possible on xv6? (Note the & in the end)

- ☐ a. no
- ☒ b. yes

The correct answer is: yes

Question 4

Complete

Mark 0.00 out of 2.00

xv6.img: bootblock kernel

```
dd if=/dev/zero of=xv6.img count=10000
dd if=bootblock of=xv6.img conv=notrunc
dd if=kernel of=xv6.img seek=1 conv=notrunc
```

Consider above lines from the Makefile. Which of the following is incorrect?

- ☒ a. The xv6.img is the virtual disk that is created by combining the bootblock and the kernel file.
- ☒ b. The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies upto 10,000 blocks on the disk.
- ☐ c. The size of xv6.img is exactly = (size of bootblock) + (size of kernel)
- ☐ d. xv6.img is the virtual processor used by the qemu emulator
- ☐ e. The size of the kernel file is nearly 5 MB
- ☐ f. Blocks in xv6.img after kernel may be all zeroes.
- ☐ g. The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies 10,000 blocks on the disk.
- ☐ h. The kernel is located at block-1 of the xv6.img
- ☐ i. The bootblock is located on block-0 of the xv6.img
- ☐ j. The bootblock may be 512 bytes or less (looking at the Makefile instruction)
- ☒ k. The size of the xv6.img is nearly 5 MB

The correct answers are: xv6.img is the virtual processor used by the qemu emulator, The xv6.img is of the size 10,000 blocks of 512 bytes each and occupies upto 10,000 blocks on the disk., The size of the kernel file is nearly 5 MB, The size of xv6.img is exactly = (size of bootblock) + (size of kernel)

Question **5**

Complete

Mark 0.43 out of 1.00

Rank the following storage systems from slowest (first) to fastest(last)

You can drag and drop the items below/above each other.

Registers

Cache

Main memory

Nonvolatile memory

Magnetic tapes

Optical disk

Hard-disk drives

The correct order for these items is as follows:

1. Magnetic tapes
2. Optical disk
3. Hard-disk drives
4. Nonvolatile memory
5. Main memory
6. Cache
7. Registers

Question **6**

Complete

Mark 1.00 out of 1.00

How does the distinction between kernel mode and user mode function as a rudimentary form of protection (security) ?

Select one:

- ☐ a. It disallows hardware interrupts when a process is running
- ☐ b. It prohibits one process from accessing other process's memory
- ☒ c. It prohibits a user mode process from running privileged instructions
- ☐ d. It prohibits invocation of kernel code completely, if a user program is running

The correct answer is: It prohibits a user mode process from running privileged instructions

Question 7

Complete

Mark 0.00 out of 2.00

Which of the following are NOT a part of job of a typical compiler?

- ☐ a. Suggest alternative pieces of code that can be written
- ☒ b. Check the program for syntactical errors
- ☐ c. Check the program for logical errors
- ☒ d. Convert high level language code to machine code
- ☒ e. Process the # directives in a C program
- ☐ f. Invoke the linker to link the function calls with their code, extern globals with their declaration

The correct answers are: Check the program for logical errors, Suggest alternative pieces of code that can be written

Question 8

Complete

Mark 0.00 out of 2.00

Match the program with it's output (ignore newlines in the output. Just focus on the count of the number of 'hi')

`main() { execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }`

hi hi

`main() { int i = fork(); if(i == 0) execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }`

hi hi hi hi

`main() { int i = NULL; fork(); printf("hi\n"); }`

No output

`main() { fork(); execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }`

hi

The correct answer is: `main() { execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }` → hi, `main() { int i = fork(); if(i == 0) execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }` → hi, `main() { int i = NULL; fork(); printf("hi\n"); }` → hi hi, `main() { fork(); execl("/usr/bin/echo", "/usr/bin/echo", "hi\n", NULL); }` → hi hi

Question **9**

Complete

Mark 0.83 out of 2.00

Select all statements that correctly explain the use/purpose of system calls.

Select one or more:

- ☒ a. Provide an environment for process creation
- ☐ b. Switch from user mode to kernel mode
- ☐ c. Handle ALL types of interrupts
- ☐ d. Handle exceptions like division by zero
- ☒ e. Run each instruction of an application program
- ☒ f. Allow I/O device access to user processes
- ☒ g. Provide services for accessing files

The correct answers are: Switch from user mode to kernel mode, Provide services for accessing files, Allow I/O device access to user processes, Provide an environment for process creation

Question **10**

Complete

Mark 0.50 out of 0.50

Compare multiprogramming with multitasking

- ☒ a. A multiprogramming system is not necessarily multitasking
- ☐ b. A multitasking system is not necessarily multiprogramming

The correct answer is: A multiprogramming system is not necessarily multitasking

Question **11**

Complete

Mark 0.60 out of 1.00

Select all the correct statements about two modes of CPU operation

Select one or more:

- ☒ a. The two modes are essential for a multitasking system
- ☐ b. Some instructions are allowed to run only in user mode, while all instructions can run in kernel mode
- ☒ c. The software interrupt instructions change the mode from user mode to kernel mode and jumps to predefined location simultaneously
- ☐ d. The two modes are essential for a multiprogramming system
- ☒ e. There is an instruction like 'iret' to return from kernel mode to user mode

The correct answers are: The two modes are essential for a multiprogramming system, The two modes are essential for a multitasking system, There is an instruction like 'iret' to return from kernel mode to user mode, The software interrupt instructions change the mode from user mode to kernel mode and jumps to predefined location simultaneously, Some instructions are allowed to run only in user mode, while all instructions can run in kernel mode

Question **12**

Complete

Mark 0.50 out of 0.50

Is the terminal a part of the kernel on GNU/Linux systems?

- ☒ a. no
- ☐ b. yes

The correct answer is: no

Question **13**

Complete

Mark 1.00 out of 1.00

Why should a program exist in memory before it starts executing ?

- ☐ a. Because the hard disk is a slow medium
- ☒ b. Because the processor can run instructions and access data only from memory
- ☐ c. Because the variables of the program are stored in memory
- ☐ d. Because the memory is volatile

The correct answer is: Because the processor can run instructions and access data only from memory

Question **14**

Complete

Mark 1.33 out of 2.00

Which of the following instructions should be privileged?

Select one or more:

- ☐ a. Read the clock.
- ☒ b. Access memory management unit of the processor
- ☐ c. Access I/O device.
- ☒ d. Turn off interrupts.
- ☐ e. Set value of a memory location
- ☒ f. Set value of timer.
- ☐ g. Access a general purpose register
- ☒ h. Modify entries in device-status table
- ☐ i. Switch from user to kernel mode.

The correct answers are: Set value of timer., Access memory management unit of the processor, Turn off interrupts., Modify entries in device-status table, Access I/O device., Switch from user to kernel mode.

Question **15**

Complete

Mark 0.07 out of 2.00

Select all the correct statements about calling convention on x86 32-bit.

- ☐ a. The ebp pointers saved on the stack constitute a chain of activation records
- ☒ b. The return value is either stored on the stack or returned in the eax register
- ☐ c. The two lines in the beginning of each function, "push %ebp; mov %esp, %ebp", create space for local variables
- ☒ d. Parameters may be passed in registers or on stack
- ☐ e. Return address is one location above the ebp
- ☐ f. Parameters may be passed in registers or on stack
- ☒ g. Compiler may allocate more memory on stack than needed
- ☐ h. Space for local variables is allocated by subtracting the stack pointer inside the code of the called function
- ☐ i. Parameters are pushed on the stack in left-right order
- ☐ j. during execution of a function, ebp is pointing to the old ebp
- ☐ k. Space for local variables is allocated by subtracting the stack pointer inside the code of the caller function

The correct answers are: Compiler may allocate more memory on stack than needed, Parameters may be passed in registers or on stack, Parameters may be passed in registers or on stack, Return address is one location above the ebp, during execution of a function, ebp is pointing to the old ebp, Space for local variables is allocated by subtracting the stack pointer inside the code of the called function, The ebp pointers saved on the stack constitute a chain of activation records

Question **16**

Complete

Mark 0.33 out of 0.50

Order the following events in boot process (from 1 onwards)

Shell	<input type="text" value="3"/>
BIOS	<input type="text" value="1"/>
Init	<input type="text" value="4"/>
OS	<input type="text" value="6"/>
Login interface	<input type="text" value="5"/>
Boot loader	<input type="text" value="2"/>

The correct answer is: Shell → 6, BIOS → 1, Init → 4, OS → 3, Login interface → 5, Boot loader → 2

[◀ \(Task\) Compulsory xv6 task](#)

Jump to...

(Optional Assignment) Shell Programming(Conformance tests) ▶

[Dashboard](#) / [My courses](#) / [Computer Engineering & IT](#) / [CEIT-Even-sem-21-22](#) / [OS-even-sem-21-22](#) / [7 February - 13 February](#)
/ [Topic-wise Quiz: 2: 9 Feb \(bootloader, memory management basics, x86\)](#)

Started on Wednesday, 9 February 2022, 7:00:12 PM

State Finished

Completed on Wednesday, 9 February 2022, 7:46:38 PM

Time taken 46 mins 26 secs

Grade 3.00 out of 11.00 (27%)

Question 1

Complete

Mark 0.00 out of 1.00

The number of GDT entries setup during boot process of xv6 is

- ☒ a. 2
- ☐ b. 3
- ☐ c. 0
- ☐ d. 256
- ☐ e. 4
- ☐ f. 255

The correct answer is: 3

Question 2

Complete

Mark 0.00 out of 1.00

x86 provides which of the following type of memory management options?

- ☐ a. segmentation and one level paging
- ☒ b. segmentation or one or two level paging
- ☐ c. segmentation and two level paging
- ☐ d. segmentation or paging
- ☐ e. segmentation and one or two level paging
- ☐ f. segmentation only

The correct answer is: segmentation and one or two level paging

Question **3**

Complete

Mark 0.00 out of 1.00

which of the following is not a difference between real mode and protected mode

- ☐ a. in real mode general purpose registers are 16 bit, in protected mode they are 32 bit
- ☐ b. in real mode the addressable memory is more than in protected mode
- ☒ c. in real mode the addressable memory is less than in protected mode
- ☐ d. in real mode the segment is multiplied by 16, in protected mode segment is used as index in GDT
- ☐ e. processor starts in real mode

The correct answer is: in real mode the addressable memory is more than in protected mode

Question **4**

Complete

Mark 0.00 out of 1.00

The kernel ELF file contains how many Program headers?

- ☐ a. 4
- ☐ b. 2
- ☐ c. 3
- ☒ d. 9
- ☐ e. 10

The correct answer is: 3

Question **5**

Not answered

Marked out of 0.50

code line, MMU setting: Match the line of xv6 code with the MMU setup employed

Answer:

The correct answer is: `inb $0x64,%al`

Question 6

Complete

Mark 1.00 out of 1.00

The kernel is loaded at Physical Address

- ☒ a. 0x00100000
- ☐ b. 0x0010000
- ☐ c. 0x80100000
- ☐ d. 0x80000000

The correct answer is: 0x00100000

Question 7

Complete

Mark 0.00 out of 1.00

Why is the code of entry() in Assembly and not in C?

- ☒ a. Because the symbol entry() is inside the ELF file
- ☐ b. Because the kernel code must begin in assembly
- ☐ c. Because it needs to setup paging
- ☐ d. There is no particular reason, it could also be in C

The correct answer is: Because it needs to setup paging

Question 8

Complete

Mark 1.00 out of 1.00

The ljmp instruction in general does

- ☐ a. change the CS and EIP to 32 bit mode, and jumps to next line of code
- ☒ b. change the CS and EIP to 32 bit mode, and jumps to new value of EIP
- ☐ c. change the CS and EIP to 32 bit mode
- ☐ d. change the CS and EIP to 32 bit mode, and jumps to kernel code

The correct answer is: change the CS and EIP to 32 bit mode, and jumps to new value of EIP

Question **9**

Complete

Mark 0.00 out of 1.00

The variable \$stack in entry.S is

- ☒ a. located at the value given by %esp as setup by bootmain()
- ☐ b. located at 0x7c00
- ☐ c. a memory region allocated as a part of entry.S
- ☐ d. located at less than 0x7c00
- ☐ e. located at 0

The correct answer is: a memory region allocated as a part of entry.S

Question **10**

Not answered

Marked out of 0.50

Match the pairs of which action is taken by whom

Answer:

The correct answer is: kernel

Question **11**

Complete

Mark 0.00 out of 1.00

ELF Magic number is

- ☐ a. 0xFFFFFFFF
- ☐ b. 0
- ☐ c. 0xELFELFELF
- ☐ d. 0xELF
- ☐ e. 0x0x464CELF
- ☒ f. 0x464C457FL
- ☐ g. 0x464C457FU

The correct answer is: 0x464C457FU

Question **12**

Complete

Mark 1.00 out of 1.00

The right side of line of code "entry = (void*)(void))(elf->entry)" means

- ☐ a. Convert the "entry" in ELF structure into void
- ☒ b. Get the "entry" in ELF structure and convert it into a function pointer accepting no arguments and returning nothing
- ☐ c. Get the "entry" in ELF structure and convert it into a function void pointer
- ☐ d. Get the "entry" in ELF structure and convert it into a void pointer

The correct answer is: Get the "entry" in ELF structure and convert it into a function pointer accepting no arguments and returning nothing

◀ Homework questions: Basics of MM, xv6 booting

Jump to...

(Code) Files, redirection, dup, (IPC)pipe ▶

[Dashboard](#) / [My courses](#) / [Computer Engineering & IT](#) / [CEIT-Even-sem-21-22](#) / [OS-even-sem-21-22](#) / [14 February - 20 February](#)
/ [Topic-wise Quiz-3 \(processes, trap handling, scheduler\)](#)

Started on Monday, 21 February 2022, 7:00:28 PM

State Finished

Completed on Monday, 21 February 2022, 7:49:24 PM

Time taken 48 mins 56 secs

Grade 8.30 out of 10.00 (83%)

Question **1**

Complete

Mark 0.80 out of 1.00

Match the elements of C program to their place in memory

Local Static variables	Data
Global variables	Data
Code of main()	Code
Function code	Code
Arguments	Stack
Mallocated Memory	Heap
#include files	No Memory needed
#define MACROS	No memory needed
Local Variables	Stack
Global Static variables	Data

The correct answer is: Local Static variables → Data, Global variables → Data, Code of main() → Code, Function code → Code, Arguments → Stack, Mallocated Memory → Heap, #include files → No memory needed, #define MACROS → No Memory needed, Local Variables → Stack, Global Static variables → Data

Question **2**

Complete

Mark 1.00 out of 1.00

What will be the output of this program

```
int main() {  
    int fd;  
    printf("%d ", open("/etc/passwd", O_RDONLY));  
    close(1);  
    fd = printf("%d ", open("/etc/passwd", O_RDONLY));  
    close(fd);  
    fd = printf("%d ", open("/etc/passwd", O_RDONLY));  
}
```

- ☐ a. 3 1 2
- ☐ b. 3 3 3
- ☒ c. 3 1 1
- ☐ d. 1 1 1
- ☐ e. 3 4 5
- ☐ f. 2 2 2

The correct answer is: 3 1 1

Question **3**

Complete

Mark 1.00 out of 1.00

Arrange in correct order, the files involved in execution of system call

vectors.S	<input type="text" value="2"/>
trap.c	<input type="text" value="4"/>
usys.S	<input type="text" value="1"/>
trapasm.S	<input type="text" value="3"/>

The correct answer is: vectors.S → 2, trap.c → 4, usys.S → 1, trapasm.S → 3

Question 4

Complete

Mark 1.00 out of 1.00

The "push 0" in vectors.S is

- ☐ a. A placeholder to match the size of struct trapframe
- ☒ b. Place for the error number value
- ☐ c. To indicate that it's a system call and not a hardware interrupt
- ☐ d. To be filled in as the return value of the system call

The correct answer is: Place for the error number value

Question 5

Complete

Mark 0.00 out of 1.00

A process blocks itself means

- ☒ a. The kernel code of an interrupt handler, moves the process to a waiting queue and calls scheduler
- ☐ b. The application code calls the scheduler
- ☐ c. The kernel code of system call, called by the process, moves the process to a waiting queue and calls scheduler
- ☐ d. The kernel code of system call calls scheduler

The correct answer is: The kernel code of system call, called by the process, moves the process to a waiting queue and calls scheduler

Question 6

Complete

Mark 1.00 out of 1.00

Select the odd one out

- ☐ a. Kernel stack of new process to Process stack of new process
- ☐ b. Process stack of running process to kernel stack of running process
- ☐ c. Kernel stack of running process to kernel stack of scheduler
- ☐ d. Kernel stack of scheduler to kernel stack of new process
- ☒ e. Kernel stack of new process to kernel stack of scheduler

The correct answer is: Kernel stack of new process to kernel stack of scheduler

Question 7

Complete

Mark 0.50 out of 0.50

Match the File descriptors to their meaning

- | | |
|---|-----------------|
| 0 | Standard Input |
| 2 | Standard error |
| 1 | Standard output |

The correct answer is: 0 → Standard Input, 2 → Standard error, 1 → Standard output

Question 8

Complete

Mark 1.00 out of 1.00

Which of the following is not a task of the code of switch() function

- ☐ a. Switch stacks
- ☒ b. Change the kernel stack location
- ☐ c. Load the new context
- ☒ d. Jump to next context EIP
- ☒ e. Save the return value of the old context code
- ☐ f. Save the old context

The correct answers are: Save the return value of the old context code, Change the kernel stack location

Question 9

Complete

Mark 0.50 out of 0.50

Match the names of PCB structures with kernel

- | | |
|-------|--------------------|
| linux | struct task_struct |
| xv6 | struct proc |

The correct answer is: linux → struct task_struct, xv6 → struct proc

Question **10**

Complete

Mark 0.50 out of 0.50

Match the MACRO with it's meaning

KERNBASE	2 GB
KERNLINK	2.224 GB
PHYSTOP	224 MB

The correct answer is: KERNBASE → 2 GB, KERNLINK → 2.224 GB, PHYSTOP → 224 MB

Question **11**

Complete

Mark 1.00 out of 1.00

The trapframe, in xv6, is built by the

- ☐ a. hardware, trapasm.S
- ☐ b. hardware, vectors.S
- ☐ c. hardware, vectors.S, trapasm.S, trap()
- ☒ d. hardware, vectors.S, trapasm.S
- ☐ e. vectors.S, trapasm.S

The correct answer is: hardware, vectors.S, trapasm.S

Question **12**

Complete

Mark 0.00 out of 0.50

Which of the following state transitions are not possible?

- ☐ a. Running -> Waiting
- ☒ b. Ready -> Waiting
- ☒ c. Waiting -> Terminated
- ☒ d. Ready -> Terminated

The correct answers are: Ready -> Terminated, Waiting -> Terminated, Ready -> Waiting

[◀ Description of some possible course mini projects](#)

Jump to...

[\(Code\) mmap related programs ▶](#)

[Dashboard](#) / [My courses](#) / [Computer Engineering & IT](#) / [CEIT-Even-sem-21-22](#) / [OS-even-sem-21-22](#) / [21 February - 27 February](#)
/ [Topic-wise Quiz-4 \(Virtual Memory\)](#)

Started on Saturday, 26 February 2022, 5:18:30 PM

State Finished

Completed on Saturday, 26 February 2022, 6:30:44 PM

Time taken 1 hour 12 mins

Grade 8.55 out of 15.00 (57%)

Question **1**

Complete

Mark 0.50 out of 0.50

Map the technique with it's feature/problem

static linking large executable file

dynamic loading allocate memory only if needed

dynamic linking small executable file

static loading wastage of physical memory

The correct answer is: static linking → large executable file, dynamic loading → allocate memory only if needed, dynamic linking → small executable file, static loading → wastage of physical memory

Question **2**

Complete

Mark 0.00 out of 1.00

Calculate the EAT in NANO-seconds (upto 2 decimal points) w.r.t. a page fault, given

Memory access time = 299 ns

Average page fault service time = 6 ms

Page fault rate = 0.8

Answer: 4.80

The correct answer is: 4800059.80

Question **3**

Complete

Mark 1.00 out of 1.00

Given six memory partitions of 300 KB , 600 KB , 350 KB , 200 KB , 750 KB , and 125 KB (in order), how would the first-fit, best-fit, and worst-fit algorithms place processes of size 115 KB and 500 KB (in order)?

best fit 115 KB	125 KB
best fit 500 KB	600 KB
worst fit 500 KB	635 KB
worst fit 115 KB	750 KB
first fit 500 KB	600 KB
first fit 115 KB	300 KB

The correct answer is: best fit 115 KB → 125 KB, best fit 500 KB → 600 KB, worst fit 500 KB → 635 KB, worst fit 115 KB → 750 KB, first fit 500 KB → 600 KB, first fit 115 KB → 300 KB

Question **4**

Complete

Mark 0.29 out of 1.00

Compare paging with demand paging and select the correct statements.

Select one or more:

- ☐ a. TLB hit ration has zero impact in effective memory access time in demand paging.
- ☒ b. Both demand paging and paging support shared memory pages.
- ☐ c. Calculations of number of bits for page number and offset are same in paging and demand paging.
- ☐ d. Demand paging requires additional hardware support, compared to paging.
- ☐ e. The meaning of valid-invalid bit in page table is different in paging and demand-paging.
- ☐ f. Paging requires NO hardware support in CPU
- ☐ g. With paging, it's possible to have user programs bigger than physical memory.
- ☐ h. With demand paging, it's possible to have user programs bigger than physical memory.
- ☐ i. Paging requires some hardware support in CPU
- ☒ j. Demand paging always increases effective memory access time.

The correct answers are: Demand paging requires additional hardware support, compared to paging., Both demand paging and paging support shared memory pages., With demand paging, it's possible to have user programs bigger than physical memory., Demand paging always increases effective memory access time., Paging requires some hardware support in CPU, Calculations of number of bits for page number and offset are same in paging and demand paging., The meaning of valid-invalid bit in page table is different in paging and demand-paging.

Question 5

Complete

Mark 0.36 out of 0.50

Map the parts of a C code to the memory regions they are related to

local variables	stack
global un-initialized variables	bss
static variables	data
global initialized variables	data
function arguments	stack
malloced memory	stack
functions	stack

The correct answer is: local variables → stack, global un-initialized variables → bss, static variables → data, global initialized variables → data, function arguments → stack, malloced memory → heap, functions → code

Question 6

Complete

Mark 0.75 out of 1.00

which of the following, do you think, are valid concerns for making the kernel pageable?

- ☒ a. The kernel must have some dedicated frames for it's own work
- ☐ b. No data structure of kernel should be pageable
- ☒ c. The kernel's own page tables should not be pageable
- ☐ d. The disk driver and disk interrupt handler should not be pageable
- ☐ e. No part of kernel code should be pageable.
- ☒ f. The page fault handler should not be pageable

The correct answers are: The kernel's own page tables should not be pageable, The page fault handler should not be pageable, The kernel must have some dedicated frames for it's own work, The disk driver and disk interrupt handler should not be pageable

Question 7

Complete

Mark 0.75 out of 1.00

W.r.t the figure given below, mark the given statements as True or False.

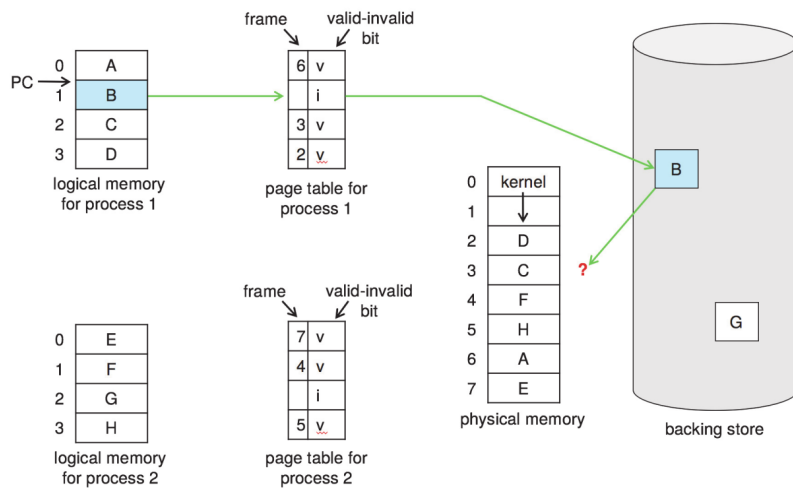


Figure 10.9 Need for page replacement.

True False

- | | | |
|----------------------------------|----------------------------------|--|
| <input checked="" type="radio"/> | <input type="radio"/> | Kernel occupies two page frames |
| <input checked="" type="radio"/> | <input type="radio"/> | Handling this scenario demands two disk I/Os |
| <input checked="" type="radio"/> | <input type="radio"/> | Local replacement means chose any of the frames 2, 3, 6 |
| <input type="radio"/> | <input checked="" type="radio"/> | Local replacement means chose any of the frame from 2 to 7 |
| <input checked="" type="radio"/> | <input type="radio"/> | Global replacement means chose any of the frame from 0 to 7 |
| <input type="radio"/> | <input checked="" type="radio"/> | Global replacement means chose any of the frame from 2 to 7 |
| <input checked="" type="radio"/> | <input type="radio"/> | The kernel's pages can not used for replacement if kernel is not pageable. |
| <input checked="" type="radio"/> | <input type="radio"/> | Page 1 of process 1 needs a replacement |

Kernel occupies two page frames: True

Handling this scenario demands two disk I/Os: True

Local replacement means chose any of the frames 2, 3, 6: True

Local replacement means chose any of the frame from 2 to 7: False

Global replacement means chose any of the frame from 0 to 7: False

Global replacement means chose any of the frame from 2 to 7: True

The kernel's pages can not used for replacement if kernel is not pageable.: True

Page 1 of process 1 needs a replacement: True

Question **8**

Complete

Mark 0.67 out of 1.00

Shared memory is possible with which of the following memory management schemes ?

Select one or more:

- ☒ a. paging
- ☒ b. segmentation
- ☐ c. continuous memory management
- ☐ d. demand paging

The correct answers are: paging, segmentation, demand paging

Question 9

Complete

Mark 0.60 out of 1.00

Given below is the "maps" file for a particular instance of "vim.basic" process.

Mark the given statements as True or False, w.r.t. the contents of the map file.

55a43501b000-55a435049000	r--p	00000000	103:05	917529	/usr/bin/vim.basic
55a435049000-55a435248000	r-xp	0002e000	103:05	917529	/usr/bin/vim.basic
55a435248000-55a4352b6000	r--p	0022d000	103:05	917529	/usr/bin/vim.basic
55a4352b7000-55a4352c5000	r--p	0029b000	103:05	917529	/usr/bin/vim.basic
55a4352c5000-55a4352e2000	rw-p	002a9000	103:05	917529	/usr/bin/vim.basic
55a4352e2000-55a4352f0000	rw-p	00000000	00:00	0	
55a436bc9000-55a436e5b000	rw-p	00000000	00:00	0	[heap]
7f275b0a3000-7f275b0a6000	r--p	00000000	103:05	917901	/usr/lib/x86_64-linux-
gnu/libnss_files-2.31.so					
7f275b0a6000-7f275b0ad000	r-xp	00003000	103:05	917901	/usr/lib/x86_64-linux-
gnu/libnss_files-2.31.so					
7f275b0ad000-7f275b0af000	r--p	0000a000	103:05	917901	/usr/lib/x86_64-linux-
gnu/libnss_files-2.31.so					
7f275b0af000-7f275b0b0000	r--p	0000b000	103:05	917901	/usr/lib/x86_64-linux-
gnu/libnss_files-2.31.so					
7f275b0b0000-7f275b0b1000	rw-p	0000c000	103:05	917901	/usr/lib/x86_64-linux-
gnu/libnss_files-2.31.so					
7f275b0b1000-7f275b0b7000	rw-p	00000000	00:00	0	
7f275b0b7000-7f275b8f5000	r--p	00000000	103:05	925247	/usr/lib/locale/locale-archive
7f275b8f5000-7f275b8fa000	rw-p	00000000	00:00	0	
7f275b8fa000-7f275b8fc000	r--p	00000000	103:05	924216	/usr/lib/x86_64-linux-
gnu/libogg.so.0.8.4					
7f275b8fc000-7f275b901000	r-xp	00002000	103:05	924216	/usr/lib/x86_64-linux-
gnu/libogg.so.0.8.4					
7f275b901000-7f275b904000	r--p	00007000	103:05	924216	/usr/lib/x86_64-linux-
gnu/libogg.so.0.8.4					
7f275b904000-7f275b905000	---p	0000a000	103:05	924216	/usr/lib/x86_64-linux-
gnu/libogg.so.0.8.4					
7f275b905000-7f275b906000	r--p	0000a000	103:05	924216	/usr/lib/x86_64-linux-
gnu/libogg.so.0.8.4					
7f275b906000-7f275b907000	rw-p	0000b000	103:05	924216	/usr/lib/x86_64-linux-
gnu/libogg.so.0.8.4					
7f275b907000-7f275b90a000	r--p	00000000	103:05	924627	/usr/lib/x86_64-linux-
gnu/libvorbis.so.0.4.8					
7f275b90a000-7f275b921000	r-xp	00003000	103:05	924627	/usr/lib/x86_64-linux-
gnu/libvorbis.so.0.4.8					
7f275b921000-7f275b932000	r--p	0001a000	103:05	924627	/usr/lib/x86_64-linux-
gnu/libvorbis.so.0.4.8					
7f275b932000-7f275b933000	---p	0002b000	103:05	924627	/usr/lib/x86_64-linux-
gnu/libvorbis.so.0.4.8					
7f275b933000-7f275b934000	r--p	0002b000	103:05	924627	/usr/lib/x86_64-linux-
gnu/libvorbis.so.0.4.8					
7f275b934000-7f275b935000	rw-p	0002c000	103:05	924627	/usr/lib/x86_64-linux-
gnu/libvorbis.so.0.4.8					
7f275b935000-7f275b937000	rw-p	00000000	00:00	0	
7f275b937000-7f275b938000	r--p	00000000	103:05	917914	/usr/lib/x86_64-linux-
gnu/libutil-2.31.so					
7f275b938000-7f275b939000	r-xp	00001000	103:05	917914	/usr/lib/x86_64-linux-
gnu/libutil-2.31.so					
7f275b939000-7f275b93a000	r--p	00002000	103:05	917914	/usr/lib/x86_64-linux-
gnu/libutil-2.31.so					
7f275b93a000-7f275b93b000	r--p	00002000	103:05	917914	/usr/lib/x86_64-linux-
gnu/libutil-2.31.so					
7f275b93b000-7f275b93c000	rw-p	00003000	103:05	917914	/usr/lib/x86_64-linux-

```

gnu/libutil-2.31.so
7f275b93c000-7f275b93e000 r--p 00000000 103:05 915906 /usr/lib/x86_64-linux-
gnu/libz.so.1.2.11
7f275b93e000-7f275b94f000 r-xp 00002000 103:05 915906 /usr/lib/x86_64-linux-
gnu/libz.so.1.2.11
7f275b94f000-7f275b955000 r--p 00013000 103:05 915906 /usr/lib/x86_64-linux-
gnu/libz.so.1.2.11
7f275b955000-7f275b956000 ---p 00019000 103:05 915906 /usr/lib/x86_64-linux-
gnu/libz.so.1.2.11
7f275b956000-7f275b957000 r--p 00019000 103:05 915906 /usr/lib/x86_64-linux-
gnu/libz.so.1.2.11
7f275b957000-7f275b958000 rw-p 0001a000 103:05 915906 /usr/lib/x86_64-linux-
gnu/libz.so.1.2.11
7f275b958000-7f275b95c000 r--p 00000000 103:05 923645 /usr/lib/x86_64-linux-
gnu/libexpat.so.1.6.11
7f275b95c000-7f275b978000 r-xp 00004000 103:05 923645 /usr/lib/x86_64-linux-
gnu/libexpat.so.1.6.11
7f275b978000-7f275b982000 r--p 00020000 103:05 923645 /usr/lib/x86_64-linux-
gnu/libexpat.so.1.6.11
7f275b982000-7f275b983000 ---p 0002a000 103:05 923645 /usr/lib/x86_64-linux-
gnu/libexpat.so.1.6.11
7f275b983000-7f275b985000 r--p 0002a000 103:05 923645 /usr/lib/x86_64-linux-
gnu/libexpat.so.1.6.11
7f275b985000-7f275b986000 rw-p 0002c000 103:05 923645 /usr/lib/x86_64-linux-
gnu/libexpat.so.1.6.11
7f275b986000-7f275b988000 r--p 00000000 103:05 924057 /usr/lib/x86_64-linux-
gnu/libltdl.so.7.3.1
7f275b988000-7f275b98d000 r-xp 00002000 103:05 924057 /usr/lib/x86_64-linux-
gnu/libltdl.so.7.3.1
7f275b98d000-7f275b98f000 r--p 00007000 103:05 924057 /usr/lib/x86_64-linux-
gnu/libltdl.so.7.3.1
7f275b98f000-7f275b990000 r--p 00008000 103:05 924057 /usr/lib/x86_64-linux-
gnu/libltdl.so.7.3.1
7f275b990000-7f275b991000 rw-p 00009000 103:05 924057 /usr/lib/x86_64-linux-
gnu/libltdl.so.7.3.1
7f275b991000-7f275b995000 r--p 00000000 103:05 921934 /usr/lib/x86_64-linux-
gnu/libtdb.so.1.4.3
7f275b995000-7f275b9a3000 r-xp 00004000 103:05 921934 /usr/lib/x86_64-linux-
gnu/libtdb.so.1.4.3
7f275b9a3000-7f275b9a9000 r--p 00012000 103:05 921934 /usr/lib/x86_64-linux-
gnu/libtdb.so.1.4.3
7f275b9a9000-7f275b9aa000 r--p 00017000 103:05 921934 /usr/lib/x86_64-linux-
gnu/libtdb.so.1.4.3
7f275b9aa000-7f275b9ab000 rw-p 00018000 103:05 921934 /usr/lib/x86_64-linux-
gnu/libtdb.so.1.4.3
7f275b9ab000-7f275b9ad000 rw-p 00000000 00:00 0
7f275b9ad000-7f275b9af000 r--p 00000000 103:05 924631 /usr/lib/x86_64-linux-
gnu/libvorbisfile.so.3.3.7
7f275b9af000-7f275b9b4000 r-xp 00002000 103:05 924631 /usr/lib/x86_64-linux-
gnu/libvorbisfile.so.3.3.7
7f275b9b4000-7f275b9b5000 r--p 00007000 103:05 924631 /usr/lib/x86_64-linux-
gnu/libvorbisfile.so.3.3.7
7f275b9b5000-7f275b9b6000 ---p 00008000 103:05 924631 /usr/lib/x86_64-linux-
gnu/libvorbisfile.so.3.3.7
7f275b9b6000-7f275b9b7000 r--p 00008000 103:05 924631 /usr/lib/x86_64-linux-
gnu/libvorbisfile.so.3.3.7
7f275b9b7000-7f275b9b8000 rw-p 00009000 103:05 924631 /usr/lib/x86_64-linux-
gnu/libvorbisfile.so.3.3.7
7f275b9b8000-7f275b9ba000 r--p 00000000 103:05 924277 /usr/lib/x86_64-linux-
gnu/libpcres2-8.so.0.9.0
7f275b9ba000-7f275ba1e000 r-xp 00002000 103:05 924277 /usr/lib/x86_64-linux-
gnu/libpcres2-8.so.0.9.0

```

```

7f275ba1e000-7f275ba46000 r--p 00066000 103:05 924277 /usr/lib/x86_64-linux-gnu/libpcr2-8.so.0.9.0
7f275ba46000-7f275ba47000 r--p 0008d000 103:05 924277 /usr/lib/x86_64-linux-gnu/libpcr2-8.so.0.9.0
7f275ba47000-7f275ba48000 rw-p 0008e000 103:05 924277 /usr/lib/x86_64-linux-gnu/libpcr2-8.so.0.9.0
7f275ba48000-7f275ba6d000 r--p 00000000 103:05 917893 /usr/lib/x86_64-linux-gnu/libc-2.31.so
7f275ba6d000-7f275bbe5000 r-xp 00025000 103:05 917893 /usr/lib/x86_64-linux-gnu/libc-2.31.so
7f275bbe5000-7f275bc2f000 r--p 0019d000 103:05 917893 /usr/lib/x86_64-linux-gnu/libc-2.31.so
7f275bc2f000-7f275bc30000 ---p 001e7000 103:05 917893 /usr/lib/x86_64-linux-gnu/libc-2.31.so
7f275bc30000-7f275bc33000 r--p 001e7000 103:05 917893 /usr/lib/x86_64-linux-gnu/libc-2.31.so
7f275bc33000-7f275bc36000 rw-p 001ea000 103:05 917893 /usr/lib/x86_64-linux-gnu/libc-2.31.so
7f275bc36000-7f275bc3a000 rw-p 00000000 00:00 0
7f275bc3a000-7f275bc41000 r--p 00000000 103:05 917906 /usr/lib/x86_64-linux-gnu/libpthread-2.31.so
7f275bc41000-7f275bc52000 r-xp 00007000 103:05 917906 /usr/lib/x86_64-linux-gnu/libpthread-2.31.so
7f275bc52000-7f275bc57000 r--p 00018000 103:05 917906 /usr/lib/x86_64-linux-gnu/libpthread-2.31.so
7f275bc57000-7f275bc58000 r--p 0001c000 103:05 917906 /usr/lib/x86_64-linux-gnu/libpthread-2.31.so
7f275bc58000-7f275bc59000 rw-p 0001d000 103:05 917906 /usr/lib/x86_64-linux-gnu/libpthread-2.31.so
7f275bc59000-7f275bc5d000 rw-p 00000000 00:00 0
7f275bc5d000-7f275bcce000 r--p 00000000 103:05 917016 /usr/lib/x86_64-linux-gnu/libpython3.8.so.1.0
7f275bcce000-7f275bf29000 r-xp 00071000 103:05 917016 /usr/lib/x86_64-linux-gnu/libpython3.8.so.1.0
7f275bf29000-7f275c142000 r--p 002cc000 103:05 917016 /usr/lib/x86_64-linux-gnu/libpython3.8.so.1.0
7f275c142000-7f275c143000 ---p 004e5000 103:05 917016 /usr/lib/x86_64-linux-gnu/libpython3.8.so.1.0
7f275c143000-7f275c149000 r--p 004e5000 103:05 917016 /usr/lib/x86_64-linux-gnu/libpython3.8.so.1.0
7f275c149000-7f275c190000 rw-p 004eb000 103:05 917016 /usr/lib/x86_64-linux-gnu/libpython3.8.so.1.0
7f275c190000-7f275c1b3000 rw-p 00000000 00:00 0
7f275c1b3000-7f275c1b4000 r--p 00000000 103:05 917894 /usr/lib/x86_64-linux-gnu/libdl-2.31.so
7f275c1b4000-7f275c1b6000 r-xp 00001000 103:05 917894 /usr/lib/x86_64-linux-gnu/libdl-2.31.so
7f275c1b6000-7f275c1b7000 r--p 00003000 103:05 917894 /usr/lib/x86_64-linux-gnu/libdl-2.31.so
7f275c1b7000-7f275c1b8000 r--p 00003000 103:05 917894 /usr/lib/x86_64-linux-gnu/libdl-2.31.so
7f275c1b8000-7f275c1b9000 rw-p 00004000 103:05 917894 /usr/lib/x86_64-linux-gnu/libdl-2.31.so
7f275c1b9000-7f275c1bb000 rw-p 00000000 00:00 0
7f275c1bb000-7f275c1c0000 r-xp 00000000 103:05 923815 /usr/lib/x86_64-linux-gnu/libgpm.so.2
7f275c1c0000-7f275c3bf000 ---p 00005000 103:05 923815 /usr/lib/x86_64-linux-gnu/libgpm.so.2
7f275c3bf000-7f275c3c0000 r--p 00004000 103:05 923815 /usr/lib/x86_64-linux-gnu/libgpm.so.2
7f275c3c0000-7f275c3c1000 rw-p 00005000 103:05 923815 /usr/lib/x86_64-linux-gnu/libgpm.so.2

```

```

7f275c3c1000-7f275c3c3000 r--p 00000000 103:05 923315 /usr/lib/x86_64-linux-gnu/libacl.so.1.1.2253
7f275c3c3000-7f275c3c8000 r-xp 00002000 103:05 923315 /usr/lib/x86_64-linux-gnu/libacl.so.1.1.2253
7f275c3c8000-7f275c3ca000 r--p 00007000 103:05 923315 /usr/lib/x86_64-linux-gnu/libacl.so.1.1.2253
7f275c3ca000-7f275c3cb000 r--p 00008000 103:05 923315 /usr/lib/x86_64-linux-gnu/libacl.so.1.1.2253
7f275c3cb000-7f275c3cc000 rw-p 00009000 103:05 923315 /usr/lib/x86_64-linux-gnu/libacl.so.1.1.2253
7f275c3cc000-7f275c3cf000 r--p 00000000 103:05 923446 /usr/lib/x86_64-linux-gnu/libcanberra.so.0.2.5
7f275c3cf000-7f275c3d9000 r-xp 00003000 103:05 923446 /usr/lib/x86_64-linux-gnu/libcanberra.so.0.2.5
7f275c3d9000-7f275c3dd000 r--p 0000d000 103:05 923446 /usr/lib/x86_64-linux-gnu/libcanberra.so.0.2.5
7f275c3dd000-7f275c3de000 r--p 00010000 103:05 923446 /usr/lib/x86_64-linux-gnu/libcanberra.so.0.2.5
7f275c3de000-7f275c3df000 rw-p 00011000 103:05 923446 /usr/lib/x86_64-linux-gnu/libcanberra.so.0.2.5
7f275c3df000-7f275c3e5000 r--p 00000000 103:05 924431 /usr/lib/x86_64-linux-gnu/libselinux.so.1
7f275c3e5000-7f275c3fe000 r-xp 00006000 103:05 924431 /usr/lib/x86_64-linux-gnu/libselinux.so.1
7f275c3fe000-7f275c405000 r--p 0001f000 103:05 924431 /usr/lib/x86_64-linux-gnu/libselinux.so.1
7f275c405000-7f275c406000 ---p 00026000 103:05 924431 /usr/lib/x86_64-linux-gnu/libselinux.so.1
7f275c406000-7f275c407000 r--p 00026000 103:05 924431 /usr/lib/x86_64-linux-gnu/libselinux.so.1
7f275c407000-7f275c408000 rw-p 00027000 103:05 924431 /usr/lib/x86_64-linux-gnu/libselinux.so.1
7f275c408000-7f275c40a000 rw-p 00000000 00:00 0
7f275c40a000-7f275c418000 r--p 00000000 103:05 924540 /usr/lib/x86_64-linux-gnu/libtinfo.so.6.2
7f275c418000-7f275c427000 r-xp 0000e000 103:05 924540 /usr/lib/x86_64-linux-gnu/libtinfo.so.6.2
7f275c427000-7f275c435000 r--p 0001d000 103:05 924540 /usr/lib/x86_64-linux-gnu/libtinfo.so.6.2
7f275c435000-7f275c439000 r--p 0002a000 103:05 924540 /usr/lib/x86_64-linux-gnu/libtinfo.so.6.2
7f275c439000-7f275c43a000 rw-p 0002e000 103:05 924540 /usr/lib/x86_64-linux-gnu/libtinfo.so.6.2
7f275c43a000-7f275c449000 r--p 00000000 103:05 917895 /usr/lib/x86_64-linux-gnu/libm-2.31.so
7f275c449000-7f275c4f0000 r-xp 0000f000 103:05 917895 /usr/lib/x86_64-linux-gnu/libm-2.31.so
7f275c4f0000-7f275c587000 r--p 000b6000 103:05 917895 /usr/lib/x86_64-linux-gnu/libm-2.31.so
7f275c587000-7f275c588000 r--p 0014c000 103:05 917895 /usr/lib/x86_64-linux-gnu/libm-2.31.so
7f275c588000-7f275c589000 rw-p 0014d000 103:05 917895 /usr/lib/x86_64-linux-gnu/libm-2.31.so
7f275c589000-7f275c58b000 rw-p 00000000 00:00 0
7f275c5ae000-7f275c5af000 r--p 00000000 103:05 917889 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7f275c5af000-7f275c5d2000 r-xp 00001000 103:05 917889 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7f275c5d2000-7f275c5da000 r--p 00024000 103:05 917889 /usr/lib/x86_64-linux-gnu/ld-2.31.so
7f275c5db000-7f275c5dc000 r--p 0002c000 103:05 917889 /usr/lib/x86_64-linux-gnu/ld-2.31.so

```

```

7f275c5dc000-7f275c5dd000 rw-p 0002d000 103:05 917889 /usr/lib/x86_64-linux-gnu/ld-
2.31.so
7f275c5dd000-7f275c5de000 rw-p 00000000 00:00 0
7ffd22d2f000-7ffd22d50000 rw-p 00000000 00:00 0 [stack]
7ffd22db0000-7ffd22db4000 r--p 00000000 00:00 0 [vvar]
7ffd22db4000-7ffd22db6000 r-xp 00000000 00:00 0 [vdso]
fffffffff600000-fffffffff601000 --xp 00000000 00:00 0 [vsyscall]

```

True	False	
<input type="radio"/>	<input checked="" type="radio"/>	The size of the heap is one page
<input checked="" type="radio"/>	<input type="radio"/>	This is a virtual memory map (not physical memory map)
<input type="radio"/>	<input checked="" type="radio"/>	vim.basic uses the math library
<input checked="" type="radio"/>	<input type="radio"/>	The 5th entry 55a4352c5000-55a4352e2000 may correspond to "data" of the vim.basic
<input checked="" type="radio"/>	<input type="radio"/>	The size of the stack is one page

The size of the heap is one page: False

This is a virtual memory map (not physical memory map): True

vim.basic uses the math library: True

The 5th entry 55a4352c5000-55a4352e2000 **may** correspond to "data" of the vim.basic: True

The size of the stack is one page: False

Question 10

Complete

Mark 0.00 out of 1.00

Select all the correct statements, w.r.t. Copy on Write

- ☐ a. Fork() used COW technique to improve performance of new process creation.
- ☒ b. Vfork() assumes that there will be no write, but rather exec()
- ☒ c. COW helps us save memory
- ☐ d. If either parent or child modifies a COW-page, then a copy of the page is made and page table entry is updated
- ☐ e. use of COW during fork() is useless if child called exit()
- ☒ f. use of COW during fork() is useless if exec() is called by the child

The correct answers are: Fork() used COW technique to improve performance of new process creation., If either parent or child modifies a COW-page, then a copy of the page is made and page table entry is updated, COW helps us save memory, Vfork() assumes that there will be no write, but rather exec()

Question **11**

Complete

Mark 1.00 out of 1.00

Assuming a 8- KB page size, what is the page numbers for the address 1093943 reference in decimal :
(give answer also in decimal)

Answer: 134

The correct answer is: 134

Question **12**

Complete

Mark 0.00 out of 1.00

Order the following events, related to page fault handling, in correct order

1. Page fault handler detects that it's a page fault and not illegal memory access
2. Disk interrupt handler runs
3. MMU detects that a page table entry is marked "invalid"
4. Page faulted process is moved to ready-queue
5. Empty frame is found
6. Page fault interrupt is generated
7. Other processes scheduled by scheduler
8. Page table of page faulted process is updated
9. Disk Interrupt occurs
10. Disk read is issued
11. Page fault handler in kernel starts executing
12. Page faulting process is made to wait in a queue

The correct order for these items is as follows:

1. MMU detects that a page table entry is marked "invalid"
2. Page fault interrupt is generated
3. Page fault handler in kernel starts executing
4. Page fault handler detects that it's a page fault and not illegal memory access
5. Empty frame is found
6. Disk read is issued
7. Page faulting process is made to wait in a queue
8. Other processes scheduled by scheduler
9. Disk Interrupt occurs
10. Disk interrupt handler runs
11. Page table of page faulted process is updated
12. Page faulted process is moved to ready-queue

Question **13**

Complete

Mark 1.00 out of 1.00

Page sizes are a power of 2 because

Select one:

- ☐ a. Certain bits are reserved for offset in logical address. Hence page size = $2^{(\text{no. of offset bits})}$
- ☐ b. Power of 2 calculations are highly efficient
- ☐ c. Certain bits are reserved for offset in logical address. Hence page size = $2^{(32 - \text{no. of offset bits})}$
- ☐ d. operating system calculations happen using power of 2
- ☐ e. MMU only understands numbers that are power of 2

The correct answer is: Certain bits are reserved for offset in logical address. Hence page size = $2^{(\text{no. of offset bits})}$

Question **14**

Complete

Mark 1.00 out of 1.00

Given below is the output of the command "ps -eo minflt,majflt,cmd" on a Linux Desktop system. Select the statements that are consistent with the output

```
626729 482768 /usr/lib/firefox/firefox -contentproc -parentBuildID 20220202182137 -prefsLen 9256 -
prefMapSize 264738 -appDir /usr/lib/firefox/browser 6094 true rdd
2167 687 /usr/sbin/apache2 -k start
1265185 222 /usr/bin/gnome-shell
102648 111 /usr/sbin/mysqld
9813 0 bash
15497 370 /usr/bin/gedit --gapplication-service
```

- ☒ a. All of the processes here exhibit some good locality of reference
- ☒ b. Firefox has likely been running for a large amount of time
- ☒ c. The bash shell is mostly busy doing work within a particular locality
- ☒ d. Apache web-server has not been doing much work

The correct answers are: Firefox has likely been running for a large amount of time, Apache web-server has not been doing much work, The bash shell is mostly busy doing work within a particular locality, All of the processes here exhibit some good locality of reference

Question **15**

Complete

Mark 0.14 out of 1.00

Suppose two processes share a library between them. The library consists of 5 pages, and these 5 pages are mapped to frames 9, 15, 23, 4, 7 respectively. Process P1 has got 6 pages, first 3 of which consist of process's own code/data and 3 correspond to library's pages 0, 2, 4. Process P2 has got 7 pages, first 3 of which consist of process's own code/data and remaining 4 correspond to library's pages 0, 1, 3, 4. Fill in the blanks for page table entries of P1 and P2.

Page table of P1, Page 5	<input type="text" value="23"/>
Page table of P1, Page 3	<input type="text" value="9"/>
Page table of P2, Page 0	<input type="text" value="6"/>
Page table of P2, Page 3	<input type="text" value="7"/>
Page table of P2, Page 4	<input type="text" value="9"/>
Page table of P2, Page 1	<input type="text" value="5"/>
Page table of P1, Page 4	<input type="text" value="9"/>

The correct answer is: Page table of P1, Page 5 → 7, Page table of P1, Page 3 → 9, Page table of P2, Page 0 → 9, Page table of P2, Page 3 → 4, Page table of P2, Page 4 → 7, Page table of P2, Page 1 → 15, Page table of P1, Page 4 → 23

Question **16**

Complete

Mark 0.50 out of 1.00

For the reference string

3 4 3 5 2

the number of page faults (including initial ones) using

FIFO replacement and 2 page frames is :

FIFO replacement and 3 page frames is :

◀ (Code) mmap related programs

Jump to...

Points from Mid-term feedback ▶

[Dashboard](#) / [My courses](#) / [Computer Engineering & IT](#) / [CEIT-Even-sem-21-22](#) / [OS-even-sem-21-22](#) / [7 March - 13 March](#)
/ [Topic-wise Quiz-5 \(xv6 memory management, userinit, exec\)](#)

Started on Monday, 7 March 2022, 7:00:12 PM

State Finished

Completed on Monday, 7 March 2022, 8:00:04 PM

Time taken 59 mins 52 secs

Grade 9.78 out of 15.00 (65%)

Question 1

Complete

Mark 1.00 out of 1.00

Why is there a call to kinit2? Why is it not merged with knit1?

- ☐ a. call to seginit() makes it possible to actually use PHYSTOP in argument to kinit2()
- ☐ b. When kinit1() is called there is a need for few page frames, but later knit2() is called to serve need of more page frames
- ☐ c. Because there is a limit on the values that the arguments to knit1() can take.
- ☒ d. knit2 refers to virtual addresses beyond 4MB, which are not mapped before kalloc() is called

The correct answer is: knit2 refers to virtual addresses beyond 4MB, which are not mapped before kalloc() is called

Question 2

Complete

Mark 1.50 out of 1.50

Arrange the following in the correct order of execution (w.r.t. 'init')

initcode() returns in forkret()

6

'initcode' process is marked RUNNABLE

3

'initcode' struct proc is created

2

userinit() is called

1

mpmain() calls scheduler()

4

initcode() calls exec("/init", ...)

8

initcode() returns from trapret()

7

scheduler() schedules initcode() process

5

The correct answer is: initcode() returns in forkret() → 6, 'initcode' process is marked RUNNABLE → 3, 'initcode' struct proc is created → 2, userinit() is called → 1, mpmain() calls scheduler() → 4, initcode() calls exec("/init", ...) → 8, initcode() returns from trapret() → 7, scheduler() schedules initcode() process → 5

Question 3

Complete

Mark 0.00 out of 2.00

exec() does this: `curproc->tf->eip = elf.entry`, but `userinit()` does this: `p->tf->eip = 0`; Select all the statements from below, that collectively explain this

- ☒ a. the 'entry' in initcode is anyways 0
- ☒ b. `exec()` loads from ELF file and the address of first instruction to be executed is given by 'entry'
- ☐ c. the initcode is created using `objcopy`, which discards all relocation information and symbols (like entry)
- ☒ d. `elf.entry` is anyways 0, so both statements mean the same
- ☒ e. In `userinit()` the function `inituvm()` has mapped the code of 'initcode' to be starting at virtual address 0
- ☐ f. the code of 'initcode' is loaded at physical address 0

The correct answers are: `exec()` loads from ELF file and the address of first instruction to be executed is given by 'entry', In `userinit()` the function `inituvm()` has mapped the code of 'initcode' to be starting at virtual address 0, the initcode is created using `objcopy`, which discards all relocation information and symbols (like entry)

Question 4

Complete

Mark 1.00 out of 1.00

What does `userinit()` do ?

- ☐ a. initializes the users
- ☐ b. sets up the 'initcode' process to start execution in `forkret()`
- ☒ c. sets up the 'initcode' process to start execution in `forkret()`
- ☐ d. sets up the 'initcode' process to start execution in `trapret()`
- ☐ e. sets up the 'init' process to start execution in `forkret()`
- ☐ f. initializes the process 'init' and starts executing it

The correct answer is: sets up the 'initcode' process to start execution in `forkret()`

Question 5

Complete

Mark 0.67 out of 1.00

Which of the following is done by mappages()?

- ☐ a. allocate page directory if required
- ☐ b. allocate page frame if required
- ☒ c. allocate page table if required
- ☒ d. create page table mappings for the range given by "va" and "va + size"
- ☐ e. create page table mappings to the range given by "pa" and "pa + size"

The correct answers are: create page table mappings for the range given by "va" and "va + size", allocate page table if required, create page table mappings to the range given by "pa" and "pa + size"

Question 6

Complete

Mark 0.00 out of 1.00

Select the statement that most correctly describes what setupkvm() does

- ☐ a. creates a 2-level page table setup with virtual->physical mappings specified in the kmap[] global array
- ☒ b. creates a 2-level page table setup with virtual->physical mappings specified in the kmap[] global array and makes kpgdir point to it
- ☐ c. creates a 1-level page table for the use by the kernel, as specified in kmap[] global array
- ☐ d. creates a 2-level page table for the use of the kernel, as specified in gdt desc

The correct answer is: creates a 2-level page table setup with virtual->physical mappings specified in the kmap[] global array

Question 7

Complete

Mark 0.00 out of 1.00

The approximate number of page frames created by kinit1 is

- ☐ a. 4
- ☐ b. 16
- ☐ c. 4000
- ☐ d. 3000
- ☐ e. 2000
- ☒ f. 1000
- ☐ g. 10

The correct answer is: 3000

Question 8

Complete

Mark 1.20 out of 1.50

Which of the following is DONE by allocproc() ?

- ☐ a. setup kernel memory mappings for the process
- ☒ b. Select an UNUSED struct proc for use
- ☐ c. ensure that the process starts in trapret()
- ☒ d. allocate kernel stack for the process
- ☐ e. allocate PID to the process
- ☒ f. ensure that the process starts in forkret()
- ☒ g. setup the trapframe and context pointers appropriately
- ☐ h. setup the contents of the trapframe of the process properly

The correct answers are: Select an UNUSED struct proc for use, allocate PID to the process, allocate kernel stack for the process, setup the trapframe and context pointers appropriately, ensure that the process starts in forkret()

Question 9

Complete

Mark 1.00 out of 1.00

Map the virtual address to physical address in xv6

KERNLINK	0x100000
0xFE000000	0xFE000000
80108000	0x108000
KERNBASE	0

The correct answer is: KERNLINK → 0x100000, 0xFE000000 → 0xFE000000, 80108000 → 0x108000, KERNBASE → 0

Question **10**

Complete

Mark 0.42 out of 1.00

Select all the correct statements about initcode

- ☐ a. code of initcode is loaded at virtual address 0
- ☒ b. The data and stack of initcode is mapped to one single page in userinit()
- ☐ c. code of 'initcode' is loaded along with the kernel during booting
- ☒ d. the size of 'initcode' is 2c
- ☐ e. code of initcode is loaded in memory by the kernel during userinit()
- ☒ f. initcode is the 'init' process
- ☒ g. initcode essentially calls exec("/init",...)

The correct answers are: code of 'initcode' is loaded along with the kernel during booting, the size of 'initcode' is 2c, The data and stack of initcode is mapped to one single page in userinit(), initcode essentially calls exec("/init",...)

Question **11**

Complete

Mark 1.00 out of 1.00

The variable 'end' used as argument to kinit1 has the value

- ☒ a. 801154a8
- ☐ b. 81000000
- ☐ c. 80110000
- ☐ d. 80102da0
- ☐ e. 8010a48c
- ☐ f. 80000000

The correct answer is: 801154a8

Question **12**

Complete

Mark 1.00 out of 1.00

What does seginit() do?

- ☐ a. Nothing significant, just repetition of earlier GDT setup but with 2-level paging setup done
- ☐ b. Nothing significant, just repetition of earlier GDT setup but with free frames list created now
- ☐ c. Adds two additional entries to GDT corresponding to Code and Data segments, but to be used in privilege level 0
- ☒ d. Adds two additional entries to GDT corresponding to Code and Data segments, but to be used in privilege level 3
- ☐ e. Nothing significant, just repetition of earlier GDT setup but with kernel page table allocated now

The correct answer is: Adds two additional entries to GDT corresponding to Code and Data segments, but to be used in privilege level 3

Question **13**

Complete

Mark 1.00 out of 1.00

Does `exec()` code around `clearptau()` lead to wastage of one page frame?

- ☐ a. no
- ☒ b. yes

The correct answer is: yes

◀ Questions for test on `kalloc/kfree/kvmalloc`, etc.

Jump to...

(Optional Assignment) Slab allocator in xv6 ▶